

Safety Not Guaranteed

Using unsafe to syscall
Windows APIs without CGO





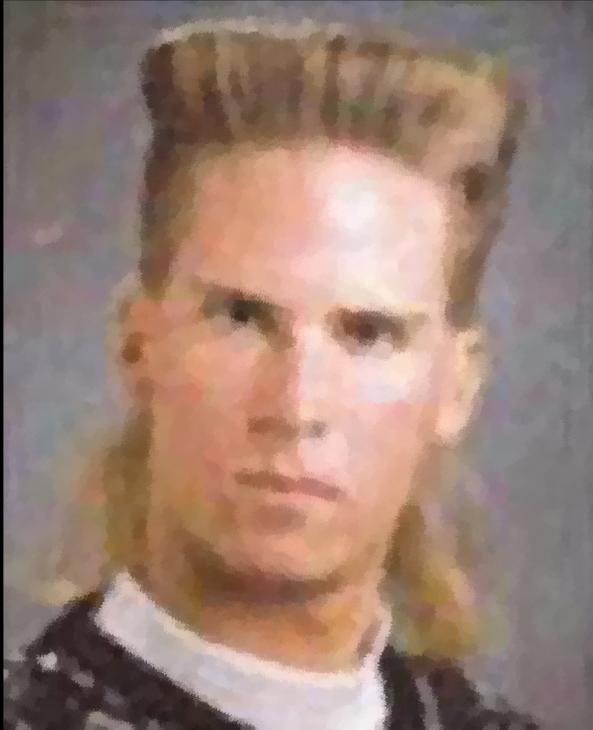
Justen Walker;
Principal Software Engineer, Global Tech
Platform



Topics

- Short intro to Syscall & unsafe.Pointer
- Loading API Functions
- Calling Windows APIs
 - Mapping C Types to/from Go Types
 - Memory Management (Managed & Unmanaged)
- Examples





WANTED: Somebody to call Win32 API functions from Go with me. This is not a joke. P.O. Box 963, New York, NY 10108. You'll learn how during this talk. Must bring your own Windows. Safety not guaranteed. I have only done this once before.



Early 2016

- K8s
- Mesos/DCOS
- Docker Swarm
- Hashicorp Nomad <- Windows!



Windows Containers

2016

Containers

- In Tech Preview (ie: not production ready)
- Big (in GB)
- Tied to OS version

Alternative: Job Objects

- Apply restrictions to EXEs
- Around since 2003
- Native to all Windows versions



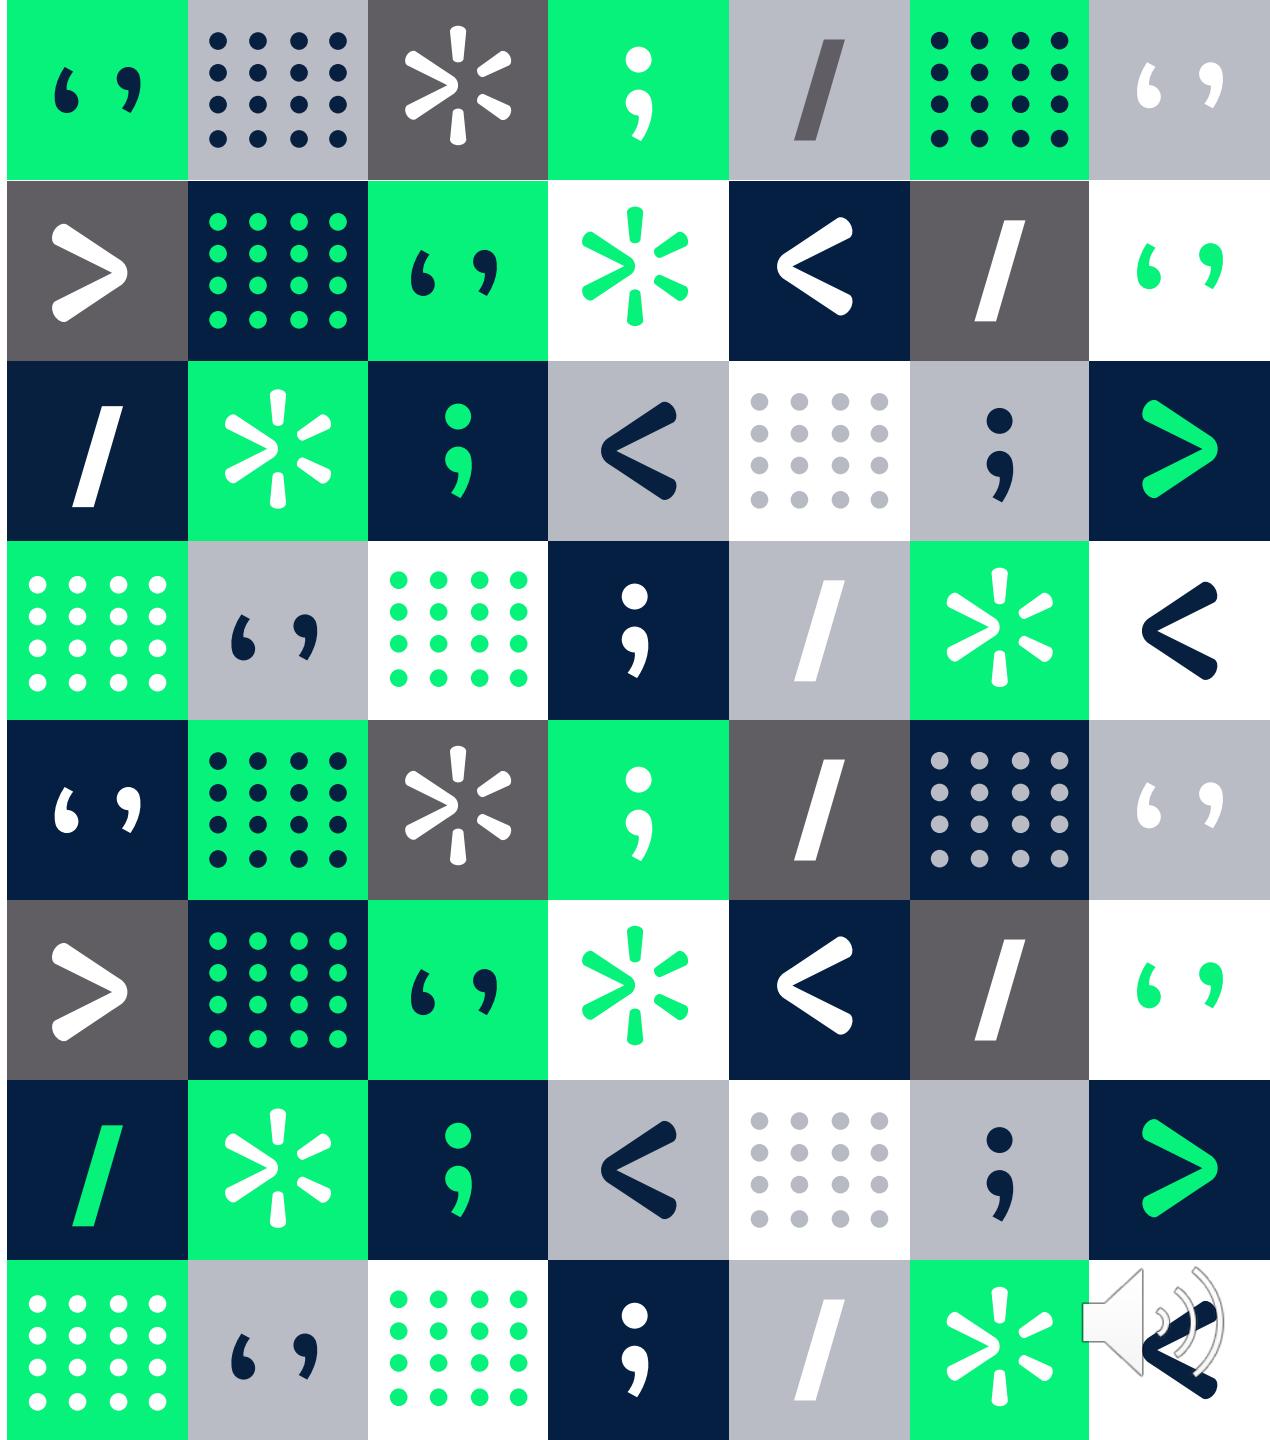
The syscall package

Your gateway to the OS



No “Syscalls” on Windows

- Microsoft does not document Syscalls
- Public API delivered via DLLs
- DLLs require C-like interface



“Syscall must always be guarded with build tags”

- // +build windows
- filename_windows.go



The unsafe package

"With the unsafe package, there are no guarantees"

- "Go Proverbs", Rob Pike 2015



unsafe.Pointer

- Type Conversions
 - $*T \leftrightarrow \text{unsafe.Pointer}$
 - $\text{unsafe.Pointer} \rightarrow \text{uintptr}$
 - $\text{uintptr} \rightarrow \text{unsafe.Pointer}$



”reinterpret_cast” of *T1 to *T2

```
type T1 struct {
    a uint32
}
type T2 struct {
    a uint16
    b uint16
}

// sizeof(T1) == sizeof(T2)

var t1 T1

t2 := *(*T2)(unsafe.Pointer(&t1))
```



Pointer Arithmetic

```
data := [...]byte{0x01,0x02,0x03}
p1 := &data[0]
// p1 = 0xc00002c001 => '0x01'

p2 := (*byte)(unsafe.Pointer(uintptr(unsafe.Pointer(&data)) + 2))
// p2 = 0xc00002c003 => '0x03'
```



Pointer Rules

- 6 Rules documented in godoc
- the compiler recognizes these
- Here's the gist
 - 1. **Never store the result of unsafe.Pointer -> uintptr in a variable**
 - 2. **Don't point to memory outside allocation boundaries**



Bad Pointers



```
data := [...]byte{0x01,0x02,0x03}

// INVALID: uintptr stored in var
just_a_number := uintptr(unsafe.Pointer(&data[0]))
p1 := (*byte)(unsafe.Pointer(just_a_number))

// INVALID: pointer past allocated boundary
p2 := (*byte)(unsafe.Pointer(uintptr(unsafe.Pointer(&data[0])) + 3))
```



Loading API Functions



Loading DLLs

Do you need to?

- golang.org/x/sys/windows
- Check for the API
- Don't do work if you don't have to



Finding DLL and Call Signatures

- docs.microsoft.com
- “Programming reference for the Win32 API”
- Example: CreateJobObject
 - CreateJobObjectA (ANSI)
 - CreateJobObjectW (UTF-16)

CreateJobObjectW function

12/05/2018 • 2 minutes to read

Creates or opens a job object.

Syntax

C++

```
HANDLE CreateJobObjectW(  
    LPSECURITY_ATTRIBUTES lpJobAttributes,  
    LPCWSTR             lpName  
)
```

Library

Kernel32.lib

DLL

Kernel32.dll



Loading a DLL



```
import "syscall"

// Loads kernel32.dll immediately
kernel32DLL = syscall.NewDLL("kernel32.dll")
// Defer loading until a procedure is called
kernel32DLL = syscall.LazyNewDLL("kernel32.dll")

import "golang.org/x/sys/windows"

// (Preferred) same as above, limited to system folders
kernel32DLL = windows.NewSystemDLL("kernel32.dll")
kernel32DLL = windows.NewLazySystemDLL("kernel32.dll")
```



Loading a DLL Procedure



```
import "golang.org/x/sys/windows"
var kernel32DLL = windows.NewLazySystemDLL("kernel32.dll")
var procCreateJobObjectW = kernel32DLL.NewProc("CreateJobObjectW")
```



Calling Windows APIs

Mapping C Types

- **Windows Data Types** reference
- Pointers are uintptr
- float/double use math package.
(See godoc for **syscall.Proc.Call**)



```
type (
    BOOL          uint32
    BOOLEAN       byte
    BYTE          byte
    DWORD         uint32
    DWORD64       uint64
    HANDLE        uintptr
    HLOCAL        uintptr
    LARGE_INTEGER int64
    LONG          int32
    LPVOID        uintptr
    SIZE_T         uintptr
    UINT          uint32
    ULONG_PTR     uintptr
    ULONGLONG     uint64
    WORD          uint16
)
```

<https://docs.microsoft.com/en-us/windows/win32/winprog/windows-data-types>



Primitive C-Type mappings

```
type (
    // 'int' is compiler dependent
    // On Windows (32 & 64bit)
    // sizeof(BOOL) = 4
    // typedef int BOOL;
    BOOL uint32
    // typedef unsigned long DWORD;
    DWORD uint32
    // typedef PVOID HANDLE;
    HANDLE uintptr
)
```

OpenProcess function

12/05/2018 • 2 minutes to read

Opens an existing local process object.

Syntax

C++

```
HANDLE OpenProcess(
    DWORD dwDesiredAccess,
    BOOL bInheritHandle,
    DWORD dwProcessId
);
```



Mapping Parameters

- Return HANDLE (uintptr)
- Args
 - *SECURITY_ATTRIBUTES
 - *uint16

CreateJobObjectW function

12/05/2018 • 2 minutes to read

Creates or opens a job object.

Syntax

C++

```
HANDLE CreateJobObjectW(
    LPSECURITY_ATTRIBUTES lpJobAttributes,
    LPCWSTR             lpName
);
```



Go string to C string

```
● ● ●  
  
// use syscall.BytePtrFromString  
func StringToCharPtr(str string) *uint8 {  
    // null terminate  
    chars := append([]byte(str), 0)  
    return &chars[0]  
}  
// use syscall.UTF16PtrFromString  
func StringToUTF16Ptr(str string) *uint16 {  
    // encode + null terminate  
    wchars := utf16.Encode([]rune(str + "\x00"))  
    return &wchars[0]  
}
```



C String to Go string

```
1 const wcharsz = uintptr(2) // uint16 = 2 bytes
2
3 func UTF16PtrToString(p *uint16) string {
4     if p == nil { return "" }
5     end := unsafe.Pointer(p)
6     var n int
7     for *(*uint16)(end) != 0 { // at NULL terminator?
8         end = unsafe.Pointer(uintptr(end) + wcharsz) // advance next char
9         n++
10    }
11    if n == 0 { return "" }
12    return string(utf16.Decode((*[1 << 30]uint16)(unsafe.Pointer(p))[:n:n]))
13 }
```



Struct Types

- Mapping is pretty much 1-to-1, Repeat for composed types

```
// typedef struct _SECURITY_ATTRIBUTES {  
type SecurityAttributes struct {  
    //    DWORD    nLength;  
    Length          uint32  
    //    LPVOID   lpSecurityDescriptor;  
    SecurityDescriptor *SECURITY_DESCRIPTOR  
    //    BOOL     bInheritHandle;  
    InheritHandle    uint32  
}  
// } SECURITY_ATTRIBUTES, *PSECURITY_ATTRIBUTES  
//     *LPSECURITY_ATTRIBUTES;
```



Calling Windows APIs – Go func to DLL call



```
func OpenProcess(desiredAccess uint32, inheritHandle bool, processId uint32) (syscall.Handle, error) {
    var bInheritHandle uint32
    if inheritHandle {
        bInheritHandle = 1
    }
    handle, _, errno := procOpenProcess.Call(
        uintptr(desiredAccess),
        uintptr(bInheritHandle),
        uintptr(processId))
    if handle == 0 {
        if errno != 0 {
            return 0, errno
        }
    }
    return syscall.Handle(handle), nil
}
```



Inspecting Return Value

Return value

If the function succeeds, the return value is an open handle to the specified process.

If the function fails, the return value is NULL. To get extended error information, call [GetLastError](#).

Remarks

...

When you are finished with the handle, be sure to close it using the [CloseHandle](#) function.



CreateJobObject

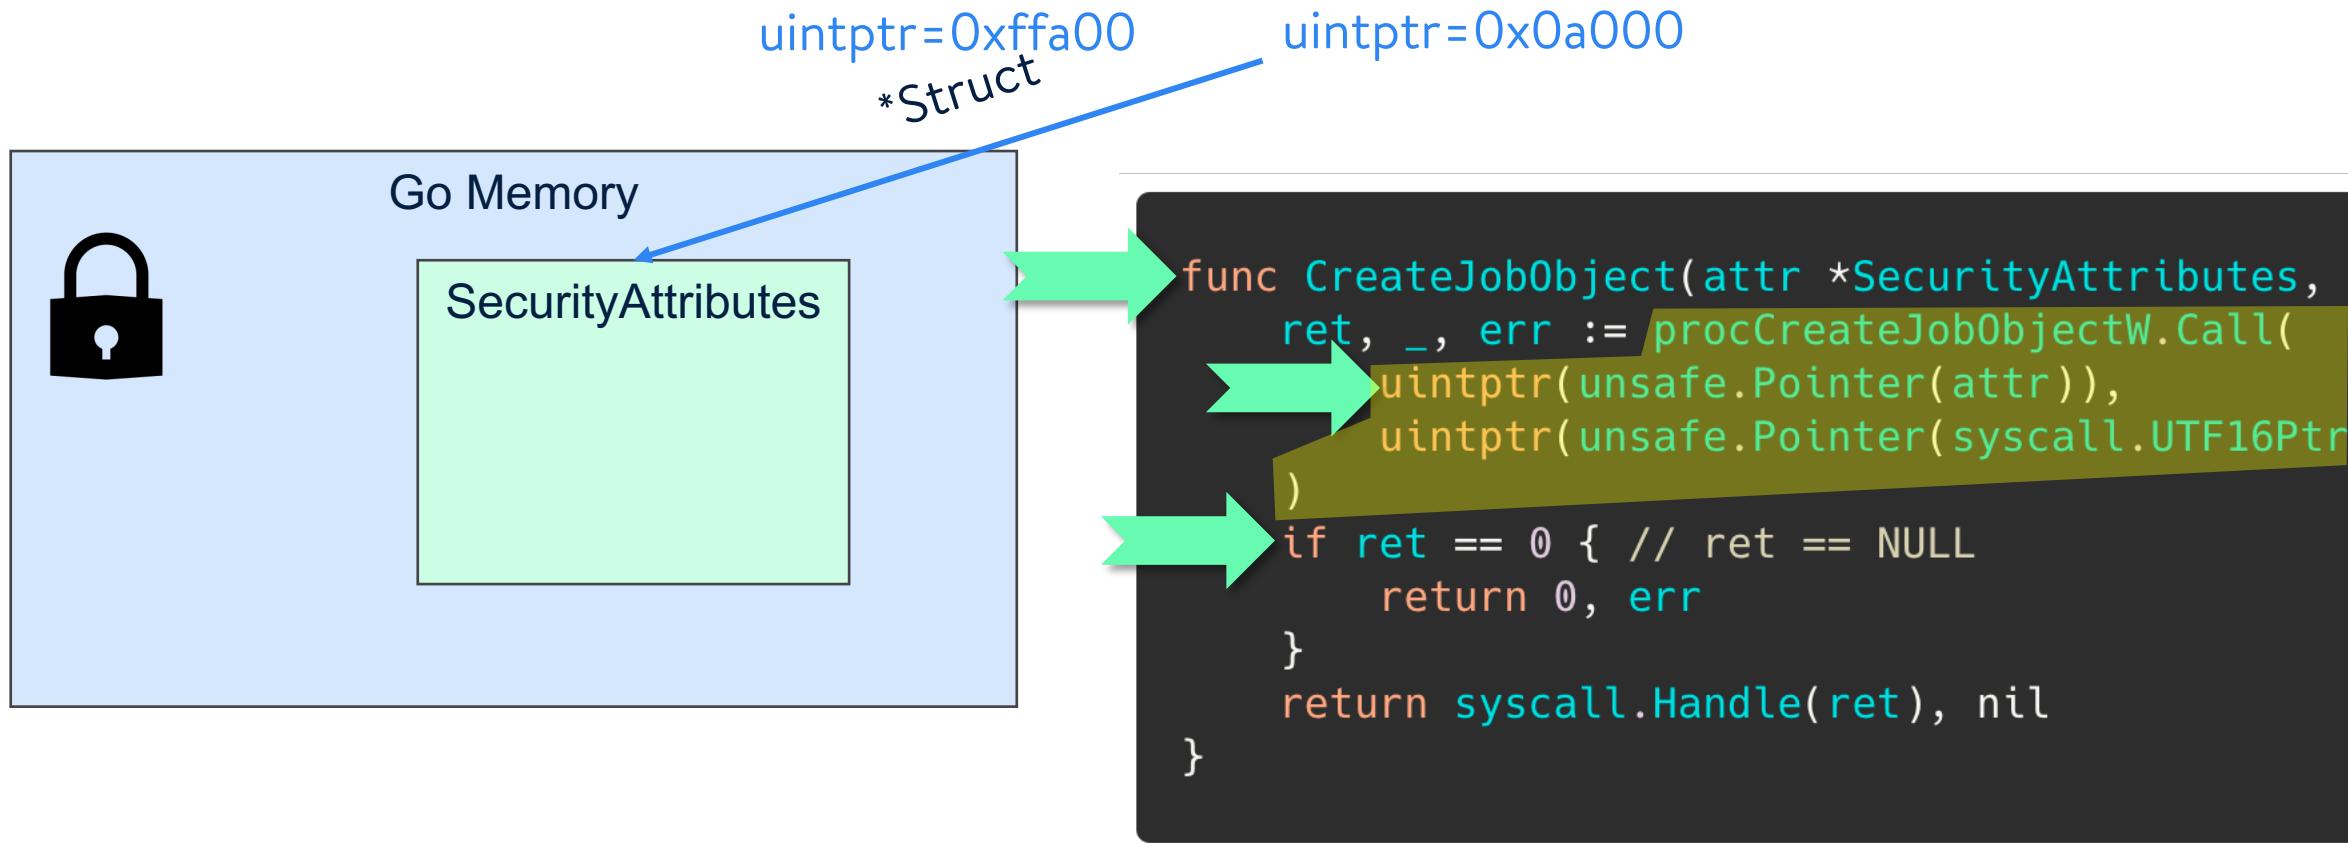
- Pass Go Pointers to C with care
- Go compiler helps here

```
func CreateJobObject(attr *SecurityAttributes, name string) (syscall.Handle, error) {
    ret, _, err := procCreateJobObjectW.Call(
        uintptr(unsafe.Pointer(attr)),
        uintptr(unsafe.Pointer(syscall.UTF16PtrFromString(name))),
    )
    if ret == 0 { // ret == NULL
        return 0, err
    }
    return syscall.Handle(ret), nil
}
```



Locking memory down during syscall

- (4) Conversion of a Pointer to a uintptr when calling syscall.Syscall



golang.org/x/sys/windows/mkwinsyscall

- Parses files for lines with //sys
- Interprets parameter lists
- Generates function wrappers
- Knows some simple conversions:
 - bool to BOOL
 - string to C-String/UTF-16 String
 - Slice to Array



mkwinsyscall - source

main.go

```
● ● ●  
package mysyscalls  
  
//go:generate go run golang.org/x/sys/windows/mkwinsyscall -output zsyscalls_windows.go  
syscalls_windows.go
```

syscalls_windows.go

```
● ● ●  
package mysyscalls  
  
//sys CreateJobObject(jobAttr *SecurityAttributes, name *uint16) (handle Handle, err error) =  
kernel32.CreateJobObjectW  
//sys GetStdHandle(stdhandle uint32) (handle Handle, err error) [failretval==InvalidHandle]
```





zsyscalls_windows.go (generated)

```
package mysyscalls

//...snip

var (
    modkernel32 = windows.NewLazySystemDLL("kernel32.dll")

    procCreateJobObjectW = modkernel32.NewProc("CreateJobObjectW")
)

func CreateJobObject(jobAttr *SecurityAttributes, name *uint16) (handle Handle, err error) {
    r0, _, e1 := syscall.Syscall(procCreateJobObjectW.Addr(), 2, uintptr(unsafe.Pointer(jobAttr)),
        uintptr(unsafe.Pointer(name)), 0)
    handle = Handle(r0)
    if handle == 0 {
        if e1 != 0 {
            err = errnoErr(e1)
        } else {
            err = syscall.EINVAL
        }
    }
    return
}
```



Memory Management

Unmanaged

Syscall allocates memory and transfers ownership to caller

```
BOOL CredEnumerateW(  
    LPCWSTR     Filter,  
    DWORD       Flags,  
    DWORD       *Count,  
    PCREDENTIALW **Credential  
) ;
```

Caller Managed

Caller allocates memory and Syscall writes to it

```
DWORD GetExtendedTcpTable(  
    PVOID          pTcpTable,  
    PDWORD         pdwSize,  
    BOOL           bOrder,  
    ULONG          ulAf,  
    TCP_TABLE_CLASS TableClass,  
    ULONG          Reserved  
) ;
```



Unmanaged Memory: CredEnumerateW

```
BOOL CredEnumerateW(  
    LPCWSTR      Filter,  
    DWORD        Flags,  
    DWORD        *Count,  
    PCREDENTIALW **Credential  
) ;
```

Pointer to an array of pointers to credentials. The returned credential is a single allocated block. Any pointers contained within the buffer are pointers to locations within this single allocated block. The single returned buffer must be freed by calling [CredFree](#).



When to Free?

- Use directly, Free when finished
- Copy and Free Immediately



Example: Designing for Unmanaged Memory

```
1 //sys CredEnumerateW(filter *uint16, flags uint32, count *uint32, credentials ***_CREDENTIALW) (err  
2   error) [failretval==0] = advapi32.CredEnumerateW  
3 //sys CredFree(buffer unsafe.Pointer) = advapi32.CredFree  
4  
5 func CredEnumerate() (*Credentials, error) {  
6     var cred Credentials  
7     var err := CredEnumerateW(nil, CRED_ENUMERATE_ALL_CREDENTIALS, &cred.count, &cred.items)  
8     if err != nil {  
9         return nil, err  
10    }  
11    free(boot)  
12    cred.items = (**_CREDENTIALW)(unsafe.Pointer(cred.items))  
13    cred.items[0].Free(boot)  
14    cred.items[0].Count = cred.count  
15    cred.items[0].Items = (**_CREDENTIALW)(unsafe.Pointer(cred.items))  
16    cred.items[0].Count = cred.count  
17 }
```

Pointer to an array of pointers to credentials. The returned credential is
a single allocated block. Any pointers contained within the buffer are



Example: Designing For Unmanaged Memory - Free

```
1 type Credentials struct {
2     free  bool
3     count uint32
4     items **_CREDENTIALW
5 }
6
7 func (c *Credentials) Free() {
8     if c.free { return }
9     CredFree(unsafe.Pointer(c.items))
10    c.free = true
11 }
```



Example: Designing For Unmanaged Memory - Iterating

```
1 func (c *Credentials) ForEach(fn func(cred Credential) error) error {
2     if c.free { return errors.New("memory freed") }
3     var sz = unsafe.Sizeof(&_amp;CREDENTIALW{})
4     for i := uint32(0); i < c.count; i++ {
5         pcred := *(**_CREDENTIALW)(unsafe.Pointer(uintptr(unsafe.Pointer(c.items)) + uintptr(i)*sz))
6         if err := fn(Credential{
7             free: &c.free,
8             pcred: pcred,
9        }); err != nil {
10            return err
11        }
12    }
13    return nil
14 }
15
16 type Credential struct {
17     free *bool
18     pcred *_CREDENTIALW
19 }
20
21 func (c *Credential) Name() string {
22     if *c.free { return "" }
23     return utf16PtrToString(c.pcred.TargetName)
24 }
25
26 // ... and other "Getters"
```

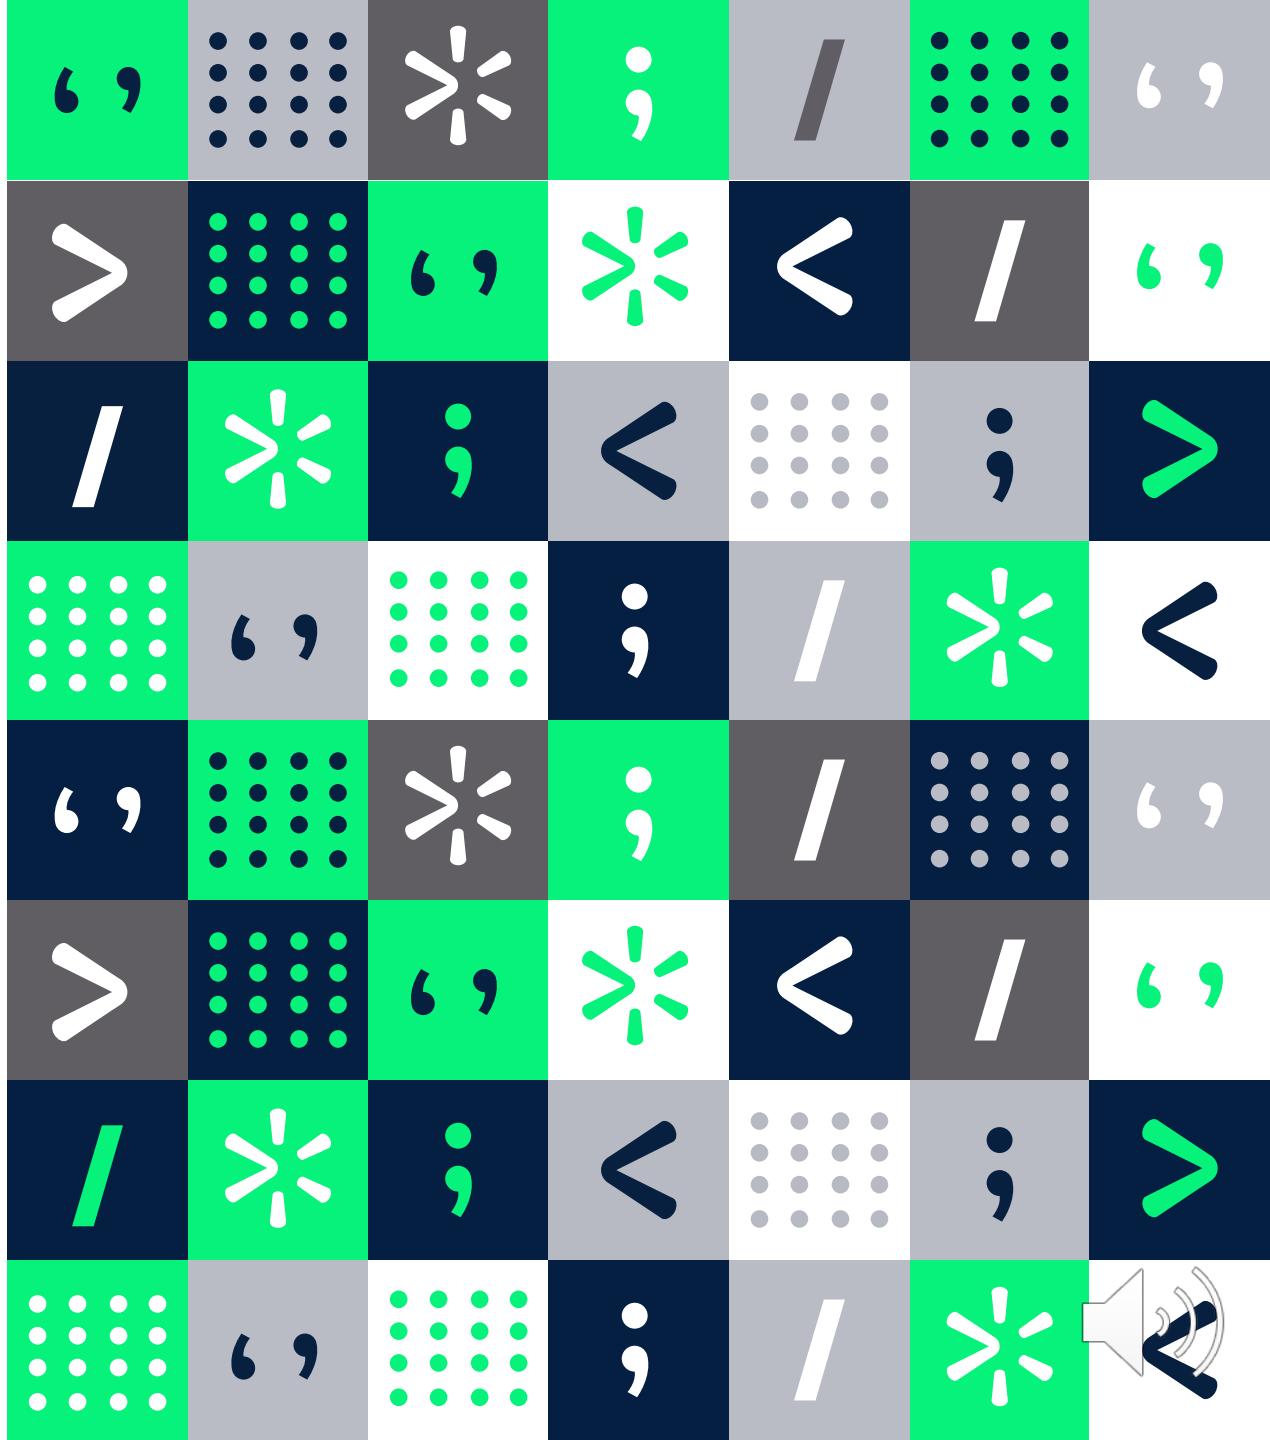


Improvement – Copy & Free

More difficult to misuse

Safer

Easier to code against



Example: Copy & Free Design - CredEnumerate

```
1 func CredEnumerate() ([]Credential, error) {
2     var count uint32
3     var items **_CREDENTIALW
4     err := CredEnumerateW(nil, CRED_ENUMERATE_ALL_CREDENTIALS, &count, &items)
5     if err != nil {
6         return nil, err
7     }
8     defer CredFree(unsafe.Pointer(items)) // Free Immediately before returning
9     var sz = unsafe.Sizeof(&_CREDENTIALW{})
10    creds := make([]Credential, 0, int(count))
11    for i := uint32(0); i < count; i++ {
12        pcred := *(**_CREDENTIALW)(unsafe.Pointer(uintptr(unsafe.Pointer(items)) + uintptr(i)*sz))
13        // toCredential: converts *_CREDENTIALW to Credential by copying
14        creds = append(creds, toCredential(pcred))
15    }
16    return creds, nil
17 }
```



Example: Copy & Free Design – Credential/Attribute

```
1 type Credential struct {
2     Name        string
3     Alias       string
4     Type        string
5     Comment     string
6     UserName    string
7     Credential  []byte
8     LastWritten time.Time
9     Attributes  []CredentialAttribute
10 }
11
12 type CredentialAttribute struct {
13     Keyword  string
14     Flags    uint32
15     Value    []byte
16 }
```



Managed Memory: Generic Example

Generic Call Pattern

```
1 var (buflen int;buffer []byte;err error;ret uintptr)
2 for {
3     ret, _, err = procSomeWindowsAPI.Call(
4         uintptr(unsafe.Pointer(&buffer[0])),
5         uintptr(unsafe.Pointer(&buflen)),
6     )
7     if err == syscall.ERROR_INSUFFICIENT_BUFFER {
8         buffer = make([]byte, buflen)
9         continue
10    }
11    break
12 }
13 result := (*WinAPIGoStruct)(unsafe.Pointer(&buffer[0])) // check ret before this
```



Managed Memory: GetExtendedTcpTable

ulAf

The version of IP used by the TCP endpoints.

Value	Meaning
AF_INET	IPv4 is used.
AF_INET6	IPv6 is used.

TableClass

The type of the TCP table structure to retrieve. This parameter can be one of the values from the [TCP_TABLE_CLASS](#) enumeration.

The [TCP_TABLE_CLASS](#) enumeration value is combined with the value of the *ulAf* parameter to determine the extended TCP information to retrieve.

```
DWORD GetExtendedTcpTable(
    PVOID pTcpTable,
    PDWORD pdwSize,
    BOOL bOrder,
    ULONG ulAf,
    TCP_TABLE_CLASS TableClass,
    ULONG Reserved
);
```



Example: GetExtendedTcpTable – Syscall Loop

```
1 //sys _GetExtendedTcpTable(tcpTable *byte, tableSize *uint32, order bool, ulAf uint32, tableClass  
2     _TCP_TABLE_CLASS, reserved uint32) (ret syscall.Errno) = iphlpapi.GetExtendedTcpTable  
3  
4 func GetTCPTable() ([]TableRow, error) {  
5     var tableSize uint32  
6     var tcpTable []byte  
7     // query for the buffer size by sending null/0  
8     ret := _GetExtendedTcpTable(nil, &tableSize, true, _AF_INET4, _TCP_TABLE_OWNER_PID_ALL, 0)  
9     for {  
10         if ret == 0 { break }  
11         if ret == _ERROR_INSUFFICIENT_BUFFER {  
12             tcpTable = make([]byte, int(tableSize))  
13             ret = _GetExtendedTcpTable(&tcpTable[0], &tableSize, true, _AF_INET4,  
14                 _TCP_TABLE_OWNER_PID_ALL, 0)  
15             continue  
16         }  
17         return nil, ret  
18     }  
19     return convertToTableRows(tcpTable), nil  
20 }
```



Example: GetExtendedTcpTable - []byte to Table

```
type TableRow struct {
    Local  *net.TCPAddr
    Remote *net.TCPAddr
    State   TcpState
    PID     int
}

func convertToTableRows(raw []byte) (result []TableRow) {
    type _MIB_TCPTABLE_OWNER_PID struct {
        NumEntries uint32
        Table      [1]_MIB_TCPROW_OWNER_PID
    }

    table := (*_MIB_TCPTABLE_OWNER_PID)(unsafe.Pointer(&raw[0]))
    n := int(table.NumEntries)
    if n == 0 { return }
    // ...
}
```

C++

```
typedef struct _MIB_TCPTABLE_OWNER_PID {
    DWORD             dwNumEntries;
    MIB_TCPROW_OWNER_PID table[ANY_SIZE];
} MIB_TCPTABLE_OWNER_PID, *PMIB_TCPTABLE_OWNER_PID;
```



Example: GetExtendedTcpTable - Table to Rows

```
// func convertToTableRows(raw []byte) (result []TableRow) { ...
    result = make([]TableRow, n)
    rows := (*[1 << 30]_MIB_TCPROW_OWNER_PID)(unsafe.Pointer(&table.Table))[0:n]
    for i, row := range rows {
        result[i] = TableRow{
            Local: convertToTCPv4Addr(row.LocalAddr, row.LocalPort),
            Remote: convertToTCPv4Addr(row.RemoteAddr, row.RemotePort),
            PID: int(row.OwningPID),
            State: TcpState(row.State),
        }
    }
    return
}
```



Full Code Available



[github.com/justenwalker/
gophercon-2020-winapi](https://github.com/justenwalker/gophercon-2020-winapi)



Thank you!());

