



CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*

# Putting an Invisible Shield on Kubernetes Secrets

Kailun Qin, Ant Group



# Agenda

- K8s Secrets: Overview
- TEE-based K8s Secrets Protection: Solution
- Production Experience @ Ant Group
- Demo
- Summary & Plan



CLOUD NATIVE + OPEN SOURCE

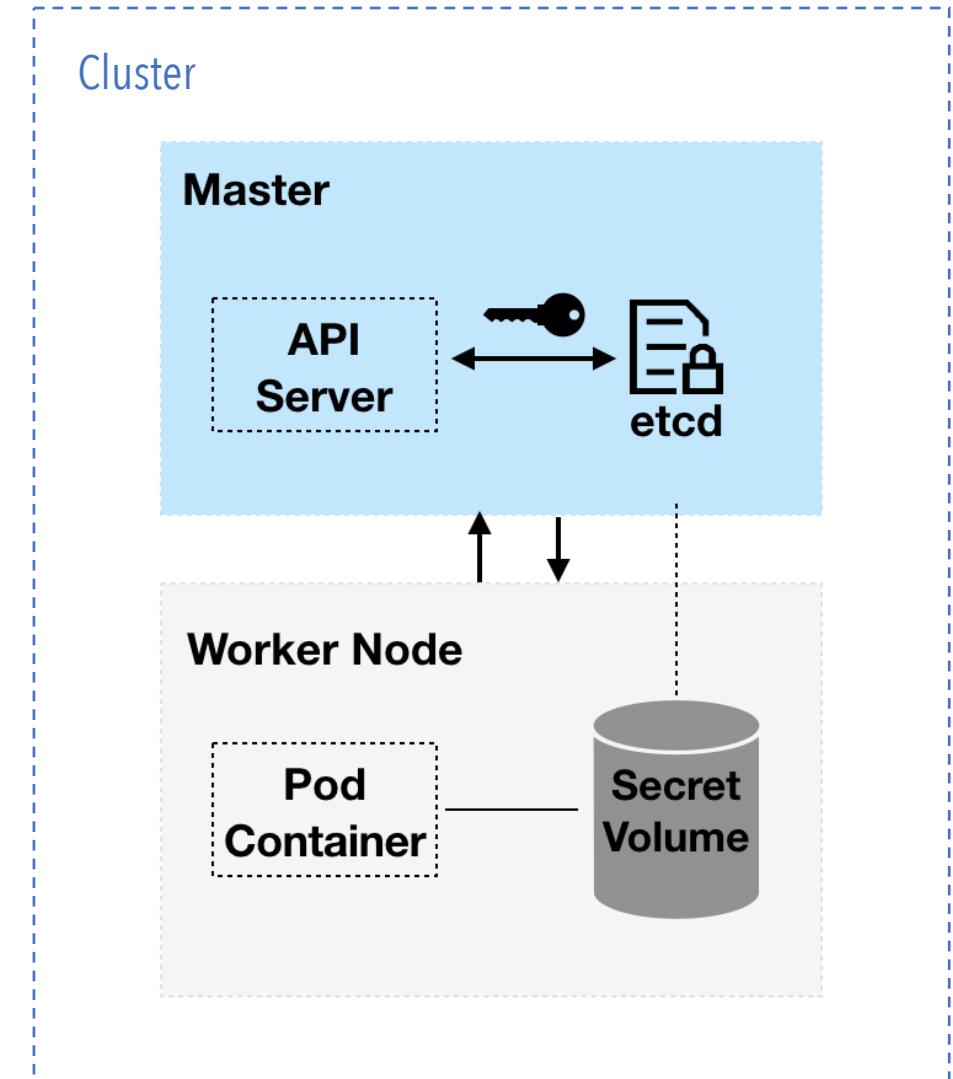
*Virtual Summit China 2020*

# K8s Secrets: Overview



# Background: K8s Secrets

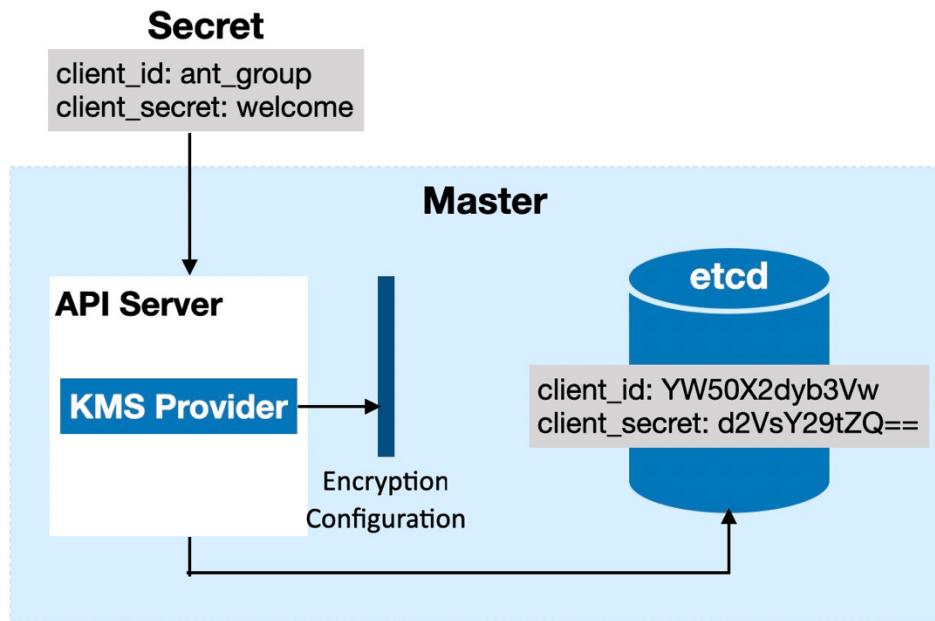
- **What they are?**
  - Sensitive information
    - Passwords
    - OAuth tokens
    - ssh keys etc.
  - Stored in etcd
    - distributed Key-Value data store
- **How about their security?**
  - Default K8s setup
    - etcd contents not encrypted (only base64 encoded)
  - > K8s 1.7+
    - at-rest encryption for etcd (local + remote)



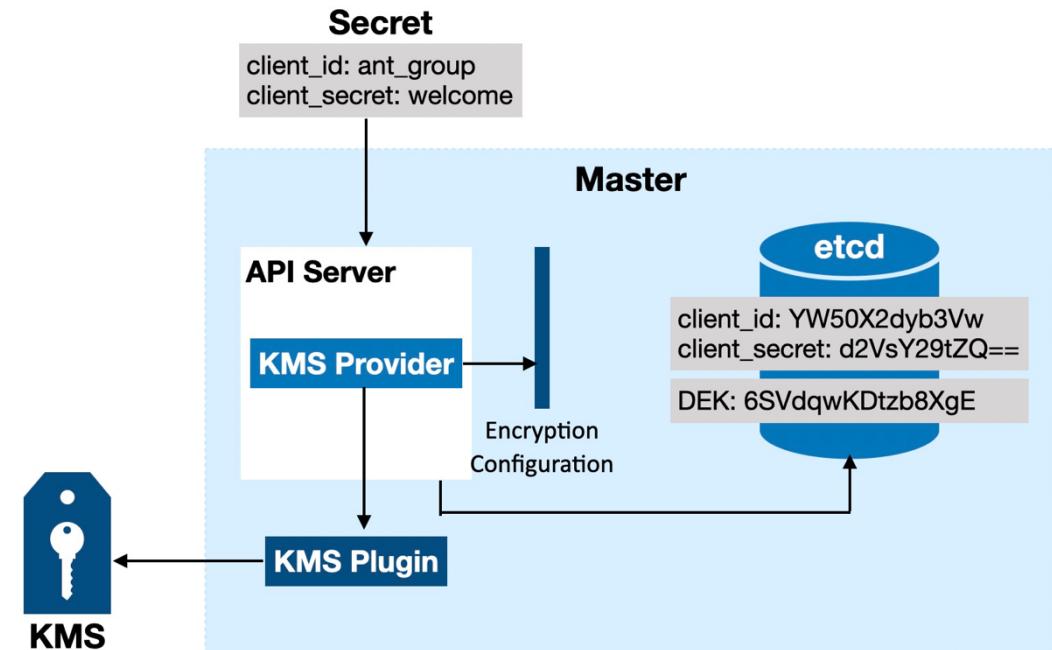


# Background: K8s Secrets

## Local Encryption Provider



## KMS Encryption Provider

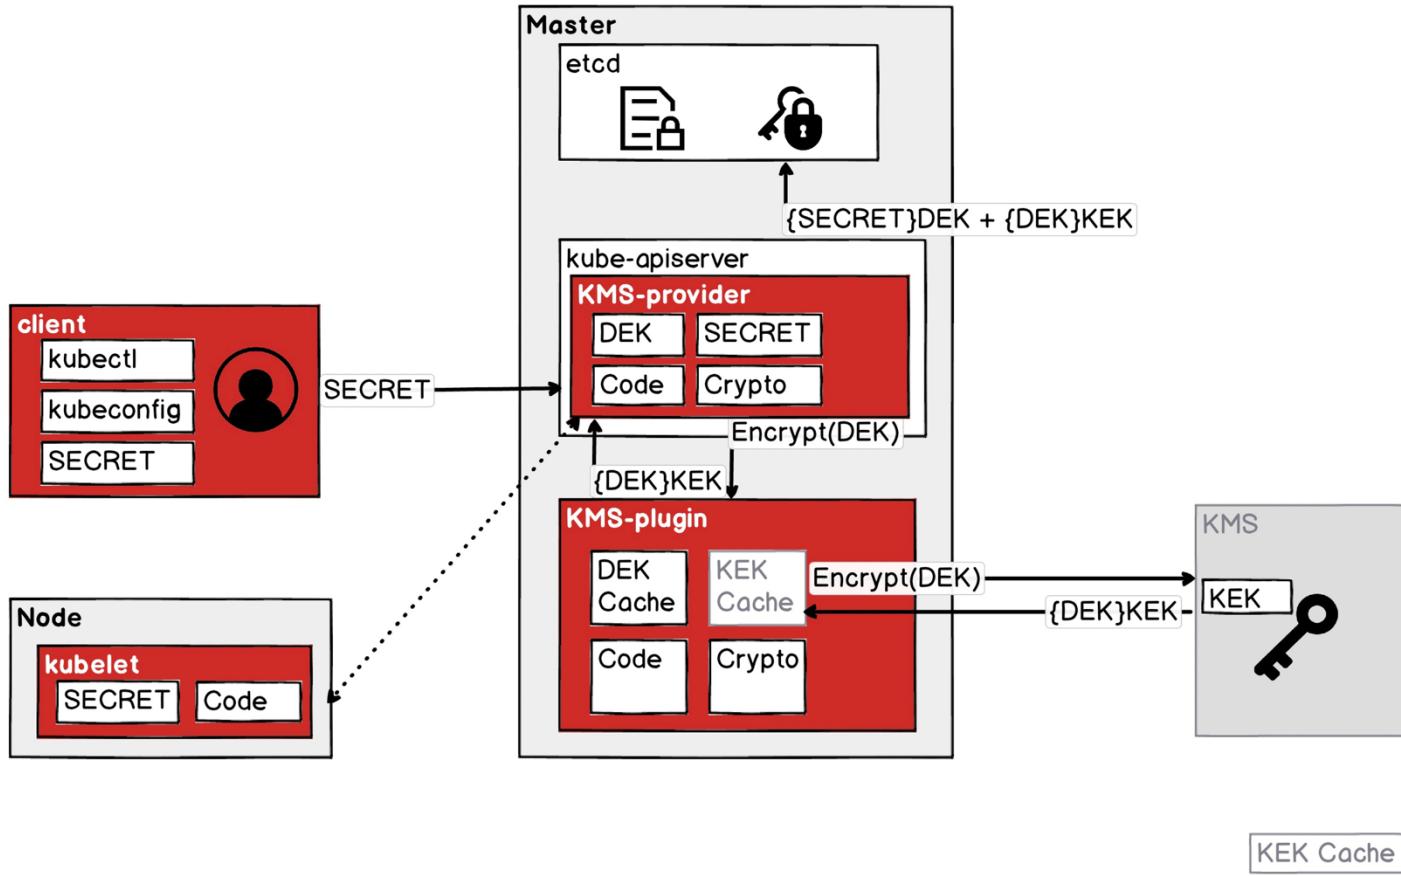


- Encryption Keys stored on API Server
- Secrets encrypted prior to storage in etcd
- Secrets decrypted on API Server prior to use

- Encryption keys stored in a remote KMS
- Use envelope encryption scheme
  - DEK & KEK



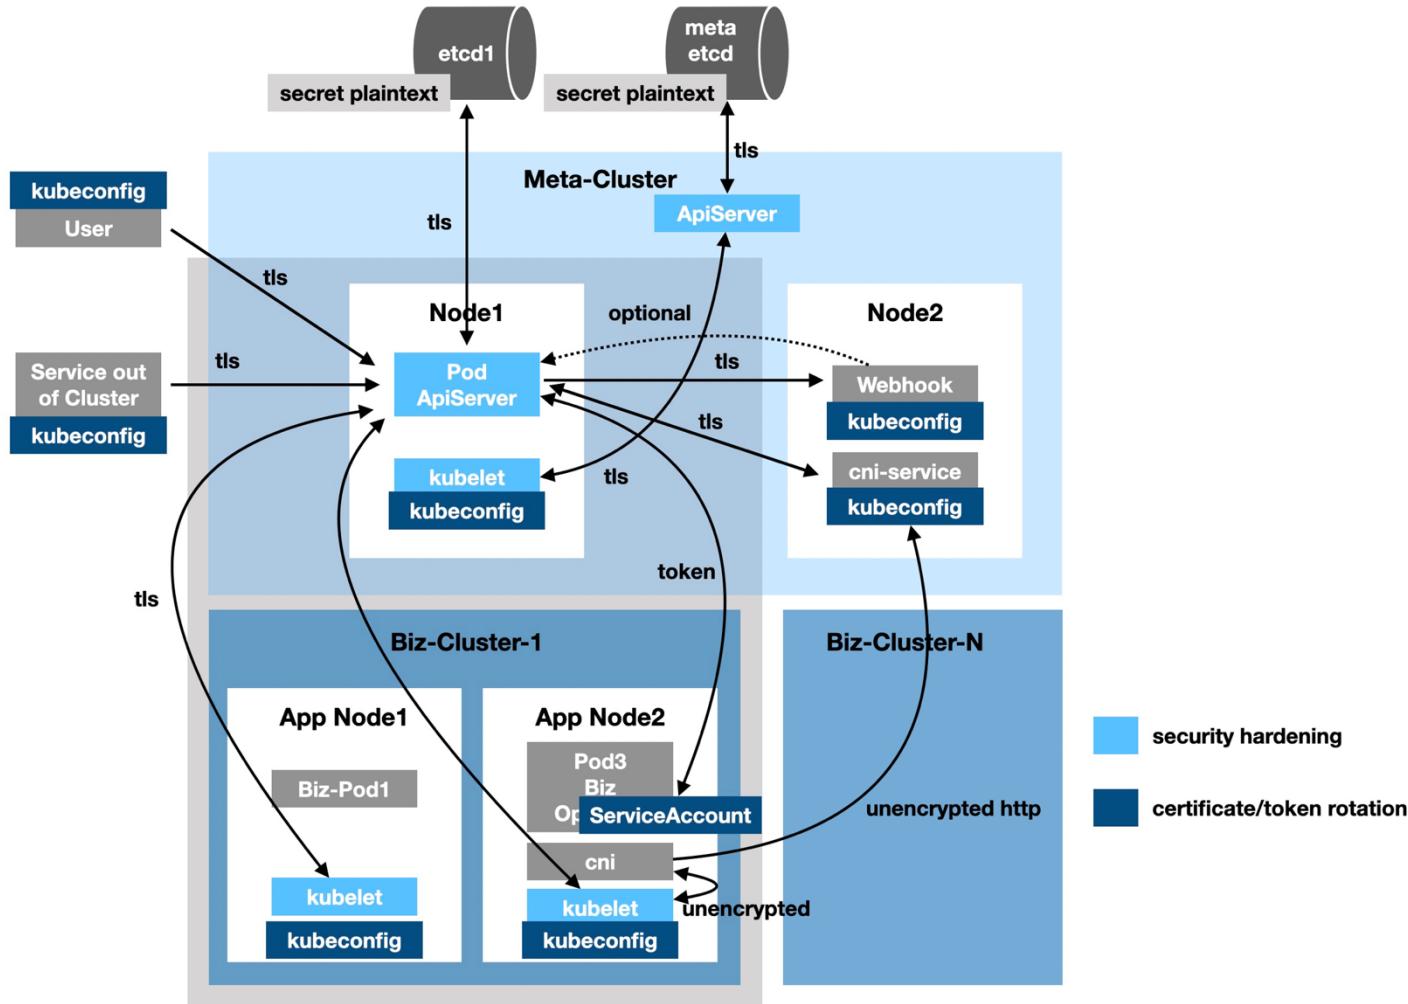
# Motivation: K8s Secrets Protection



- **Performance & latency**
  - Network
- **Security**
  - DEK in the clear in memory
  - Secret in the clear in memory
  - kubeconfig in the clear in memory
  - KEK in the clear in memory
    - ✓ leak ALL DEKs
    - ✓ leak ALL secrets
    - ✓ **trust collapse!**
  - DEK decryption interfaces invoked by fake users



# Motivation: K8s Secrets Protection



- Kube-on-Kube [1]
  - ✓ Components => **too many!**
  - ✓ Interactions => **complicated!**
  - ✓ User access management => **raw and extensive!**
  - ✓ Secrets management => **crucial!**
- Financial-grade security



CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*

# TEE-based Secrets Protection: Solution



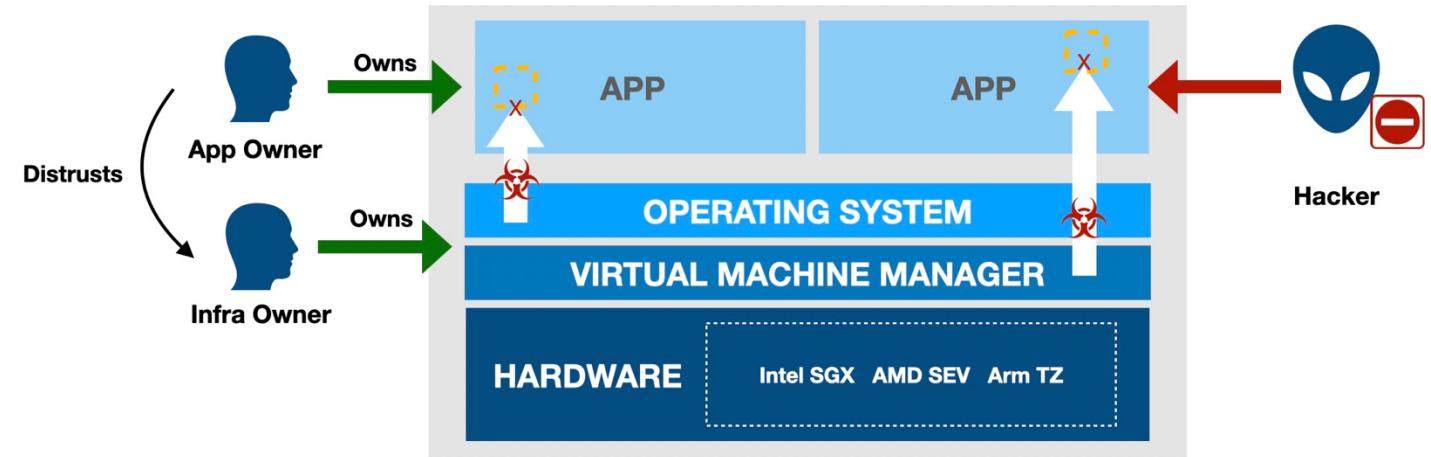
# Confidential Computing

A **Trusted Execution Environment** (TEE) is

- a secure area protected by the processor (aka. Enclave)

Example: Intel SGX

- Strong isolation
- Encrypted memory
- SW/HW attacks prevented



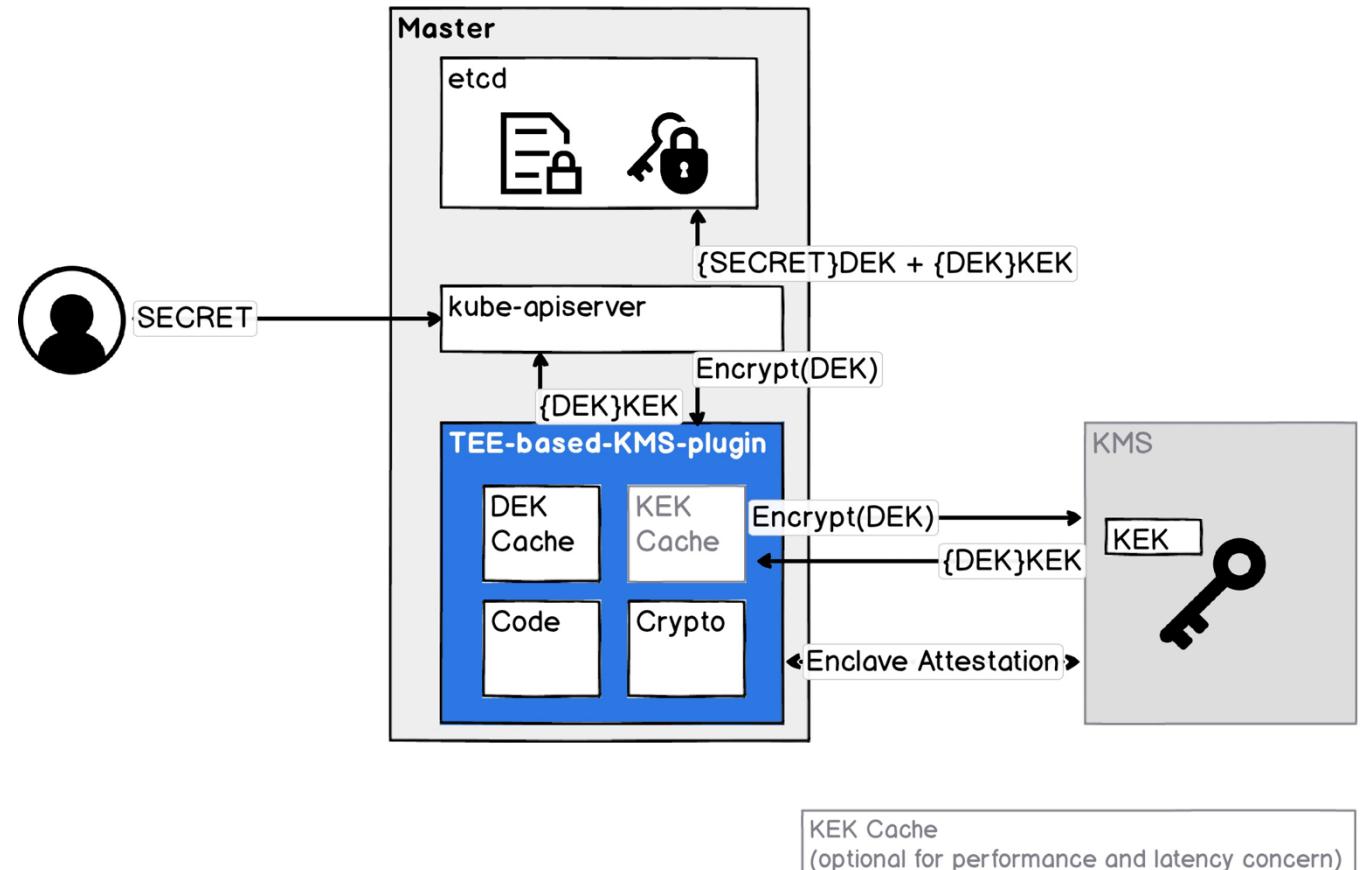


CLOUD NATIVE + OPEN SOURCE

Virtual Summit China 2020

# TEE-based KMS Plugin [1]

- Address performance & latency concerns
  - Reduce / minimize remote KMS interactions w/o compromising security
- Address security threats
  - etcd compromise
  - Host (KMS plugin) compromise
    - leak DEKs
    - leak KEKs



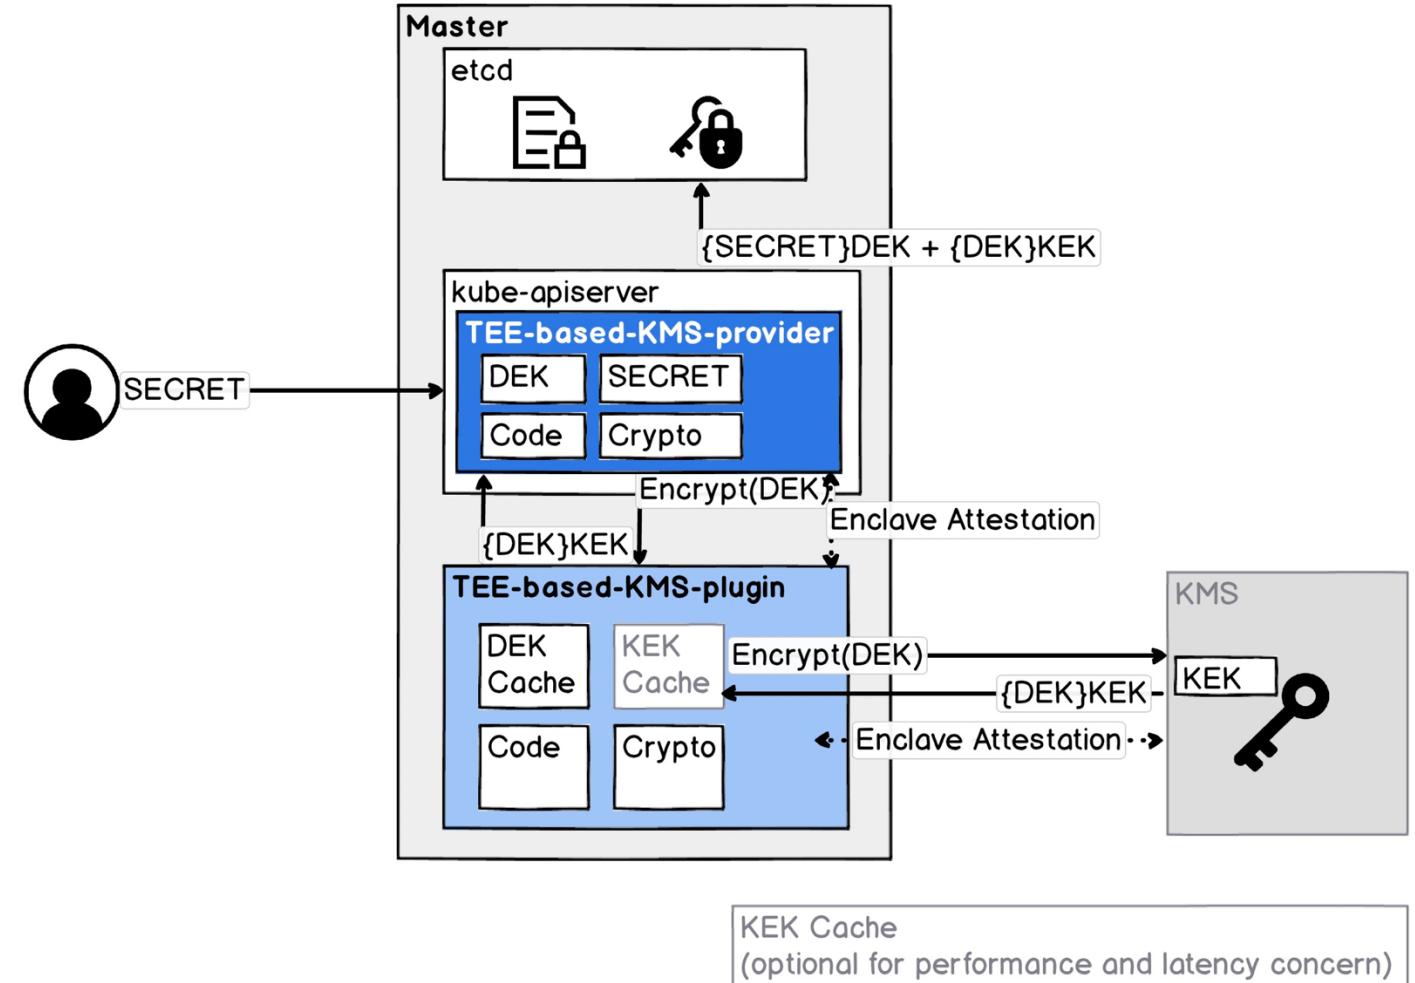


CLOUD NATIVE + OPEN SOURCE

Virtual Summit China 2020

# TEE-based KMS Provider

- Address security threats
  - Host(KMS provider) compromise
    - leak DEKs
    - leak Secrets
  - Fraudsters calling DEK decryption interfaces



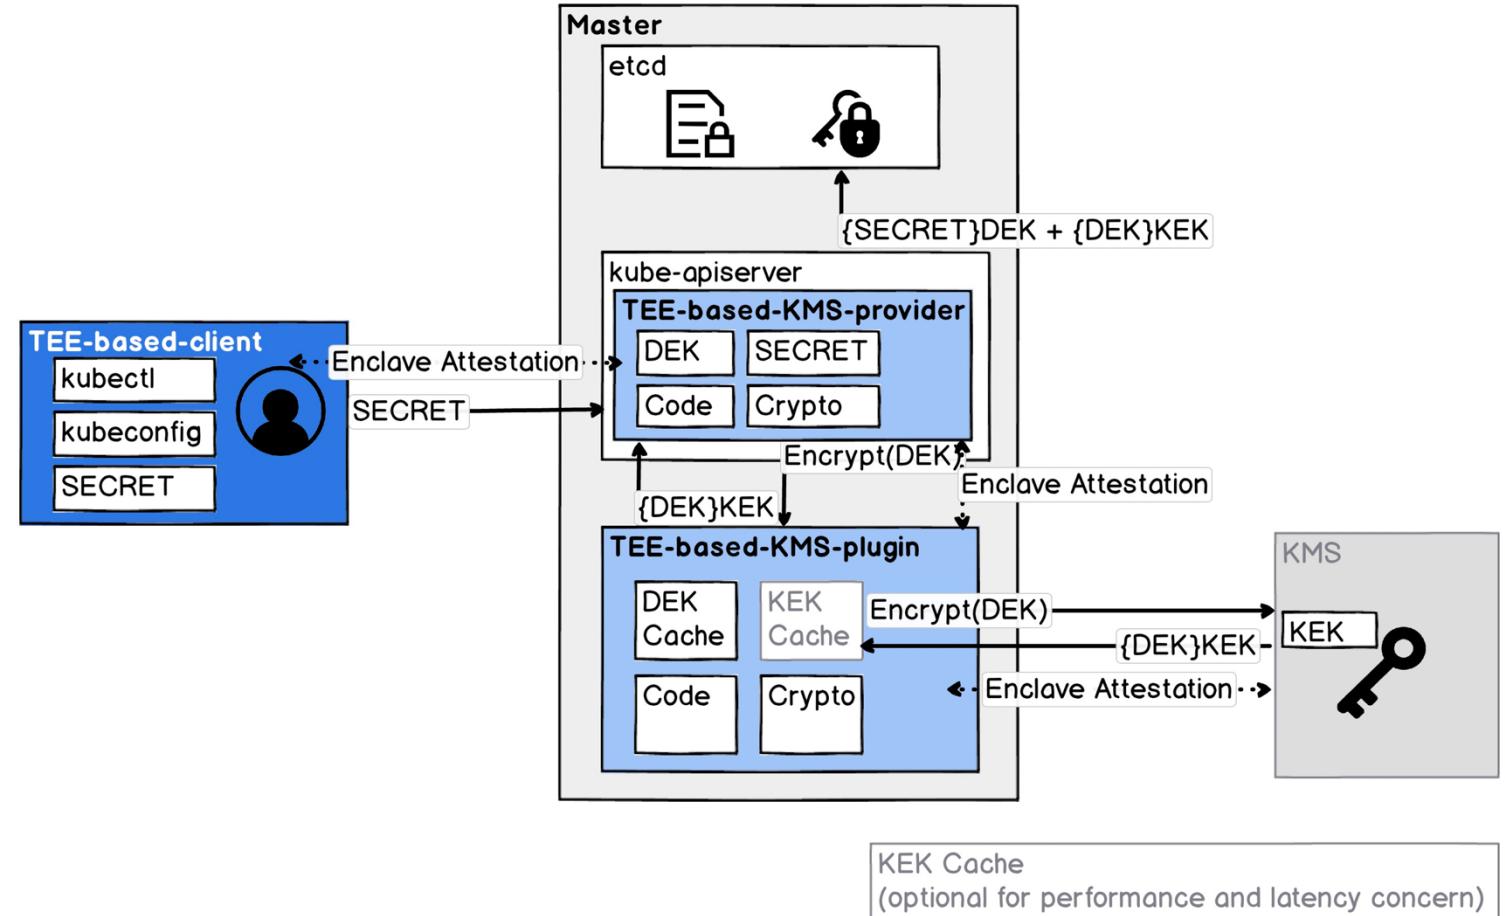


CLOUD NATIVE + OPEN SOURCE

Virtual Summit China 2020

# TEE-based Kubectl

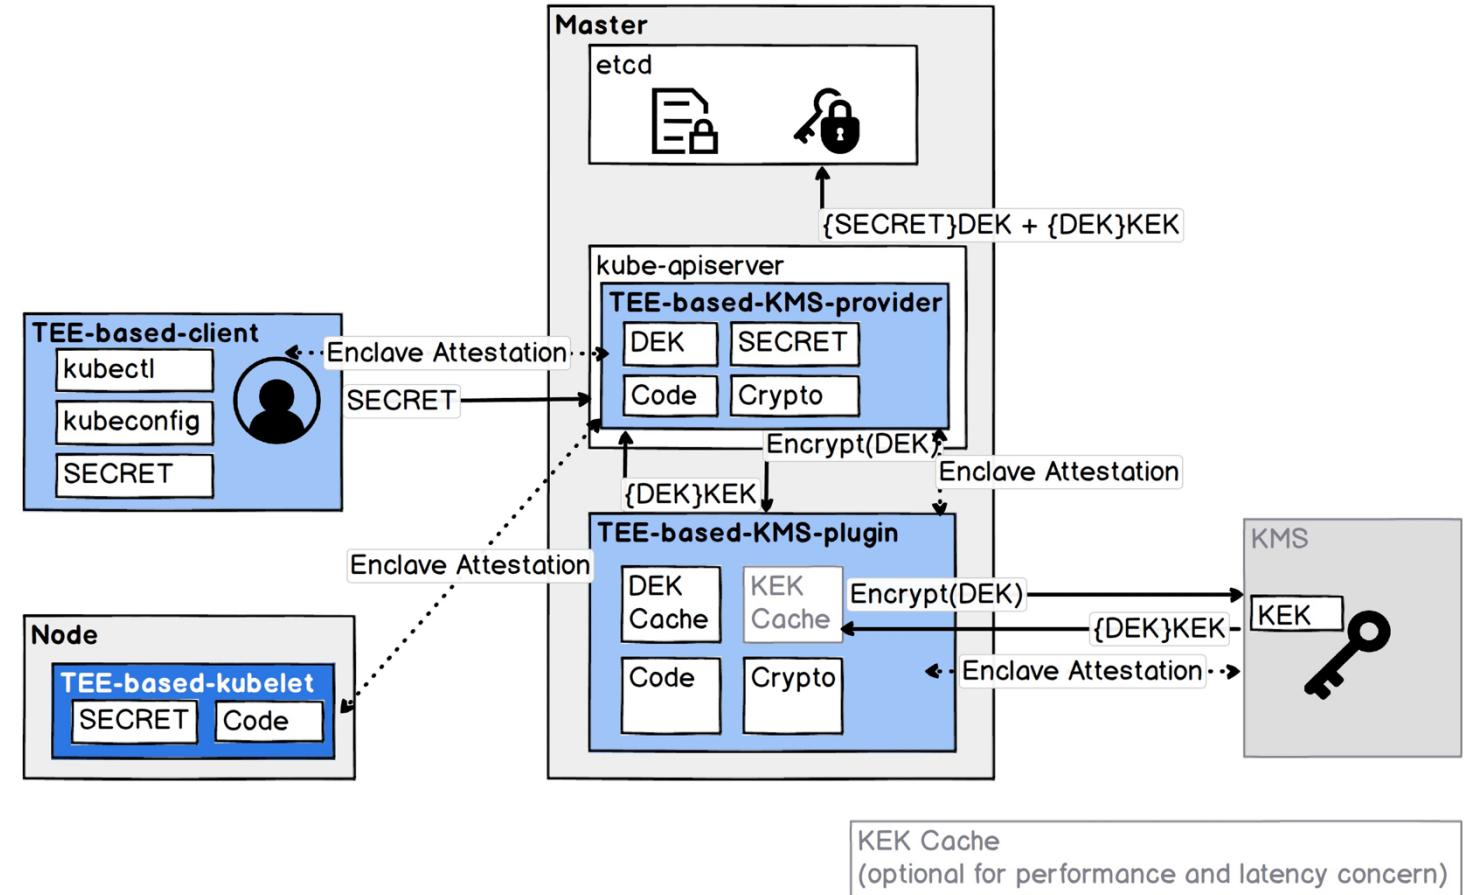
- Address security threats
  - Client compromise
    - kubeconfig maliciously reused by attackers
    - kubeconfig in the clear in clients' memory
    - leak users' secrets
  - Sending to / receiving from malicious software entity (logic)





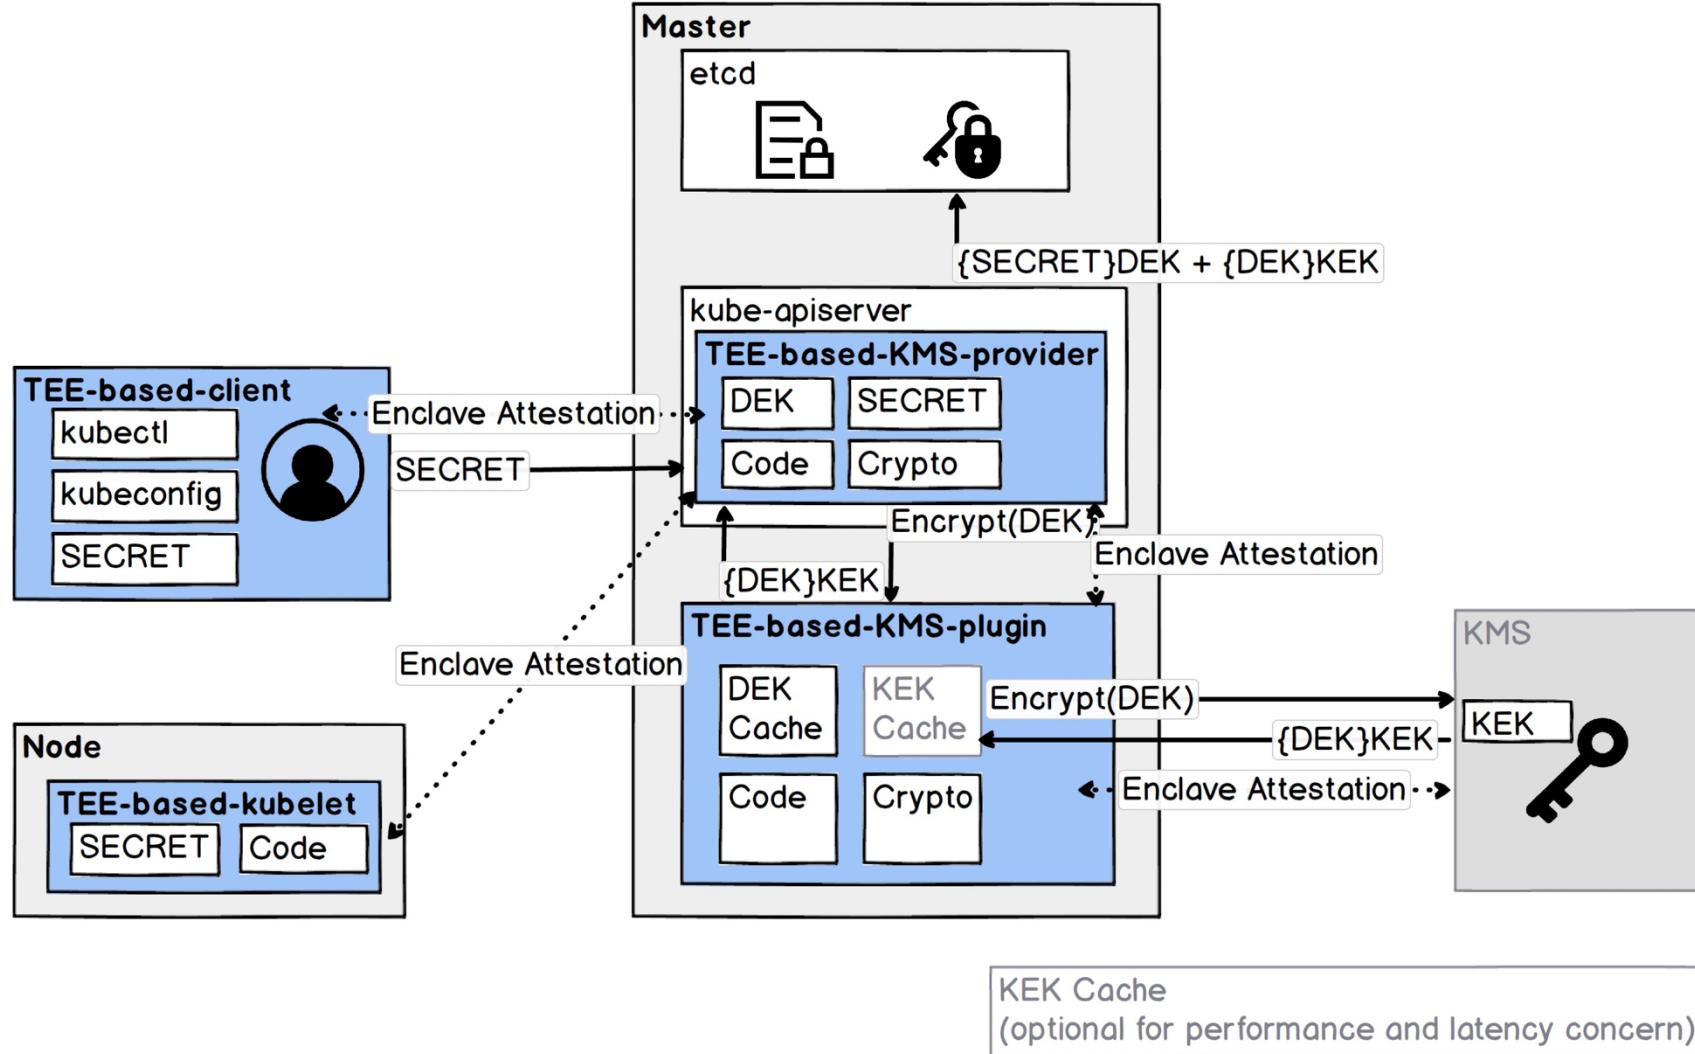
# TEE-based Kubelet

- Address security threats
  - Node (kubelet) compromise
    - leak secrets on consumption



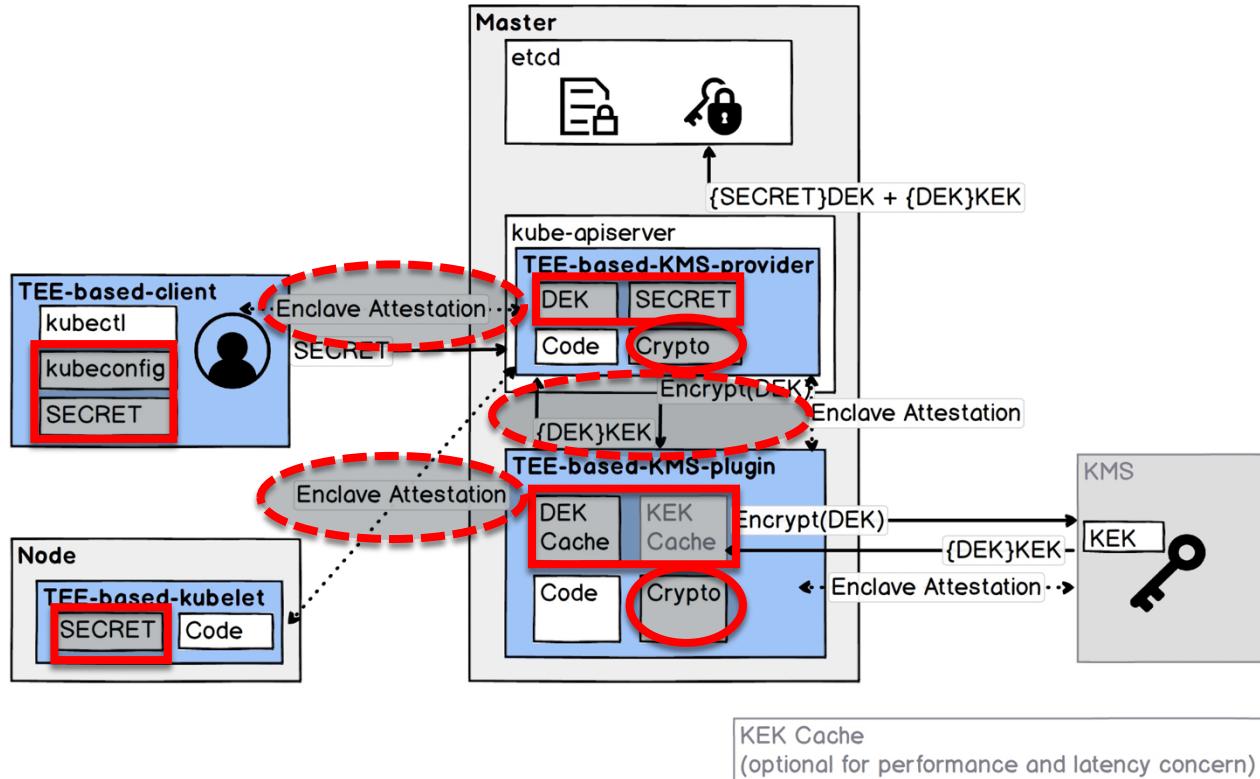


# TEE-based Secrets Protection





# TEE-based Secrets Protection (cont.)



## In a nutshell...

- Protecting DEKs / KEKs / Secrets / kubeconfig in memory
- Protecting the encryption / decryption of DEKs / Secrets
- Introducing mutual (remote / local) attestations between entities



CLOUD NATIVE + OPEN SOURCE

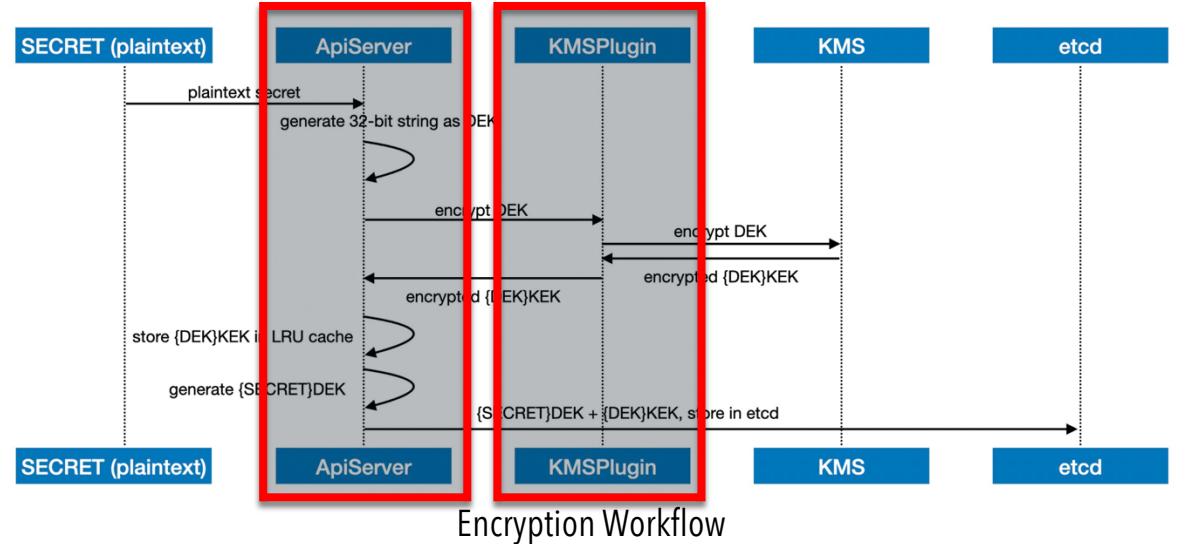
*Virtual Summit China 2020*

# Production Experience @ Ant Group

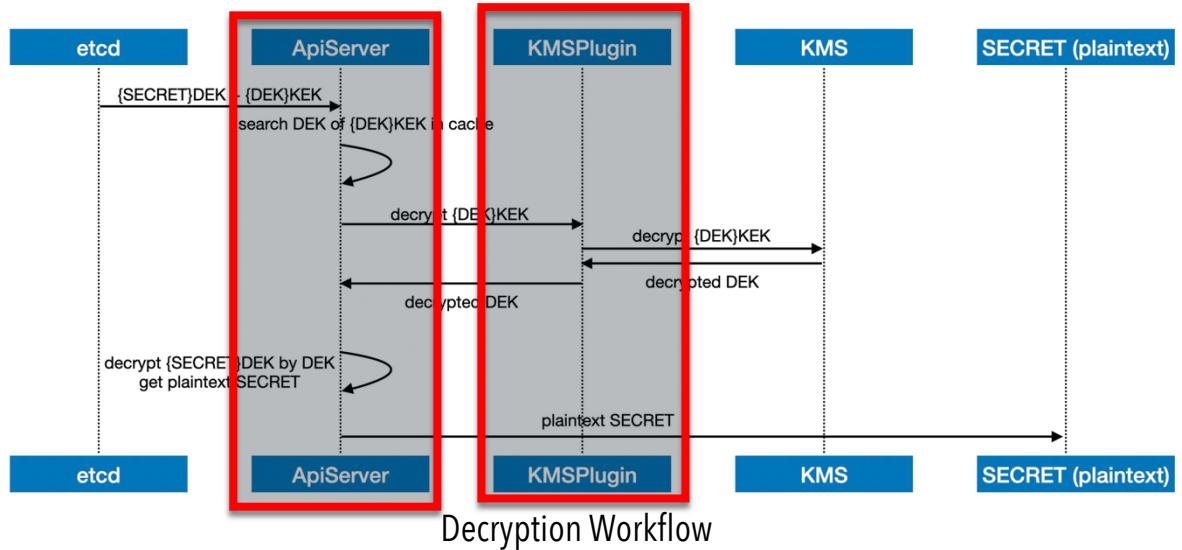


# KMS Plugin

- **Workflow**
  - Encryption
  - Decryption
- **Engineering decisions**
  - apiserver is responsible for
    - DEK generation
    - Secret en/decryption
  - kms-plugin
    - keeps KEK cache
    - only en/decrypts DEK, not secrets



Encryption Workflow



Decryption Workflow



# KMS Plugin (cont.)

- **Deployment Modes**

- One kms-plugin container per Master Node: sidecar to apiserver
- Use Annotation to enable encrypted secret read / write
- LivenessProbe for health check

app config  
identity credentials to KMS  
static keys

kms-plugin config

- **Configurations**

- kms-plugin
- apiserver

- **Caching**

- API server
  - Set up Encrypted(DEK) => DEK mapping
- KMS plugin
  - Set up SecretKeyName:SecretKeyVersion => SecretKeyData mapping

```
1 apiVersion: apiserver.config.k8s.io/v1
2 kind: EncryptionConfig
3 resources:
4   - resources:
5     - secrets
6   providers:
7     - kms:
8       name: kms-plugin
9       endpoint: unix:///var/run/kmsplugin/kms-plugin.sock
10      cachesize: 100
11      timeout: 3s
12      - identity: {}
```

apiserver config



# KMS Plugin (cont.)

- **Gray release**
  - extensions-webhook: /mutating-secret
  - Annotation: /storage-transform-disable=<bool>
- **Emergency management**
  - High Availability guarantee
    - KMS
    - API server & kms-plugin
  - Cron job backup for KEKs (from KMS)
  - Static key configuration support in kms-plugin
  - One click decryption
  - Key force update
  - Liveness probe
- **Monitoring**
  - Integration w/ Prometheus
  - Metrics including
    - latency of en/decryption
    - failure times of en/decryption
    - KMS health check
- **Ops tooling**
  - kms-plugin-tools



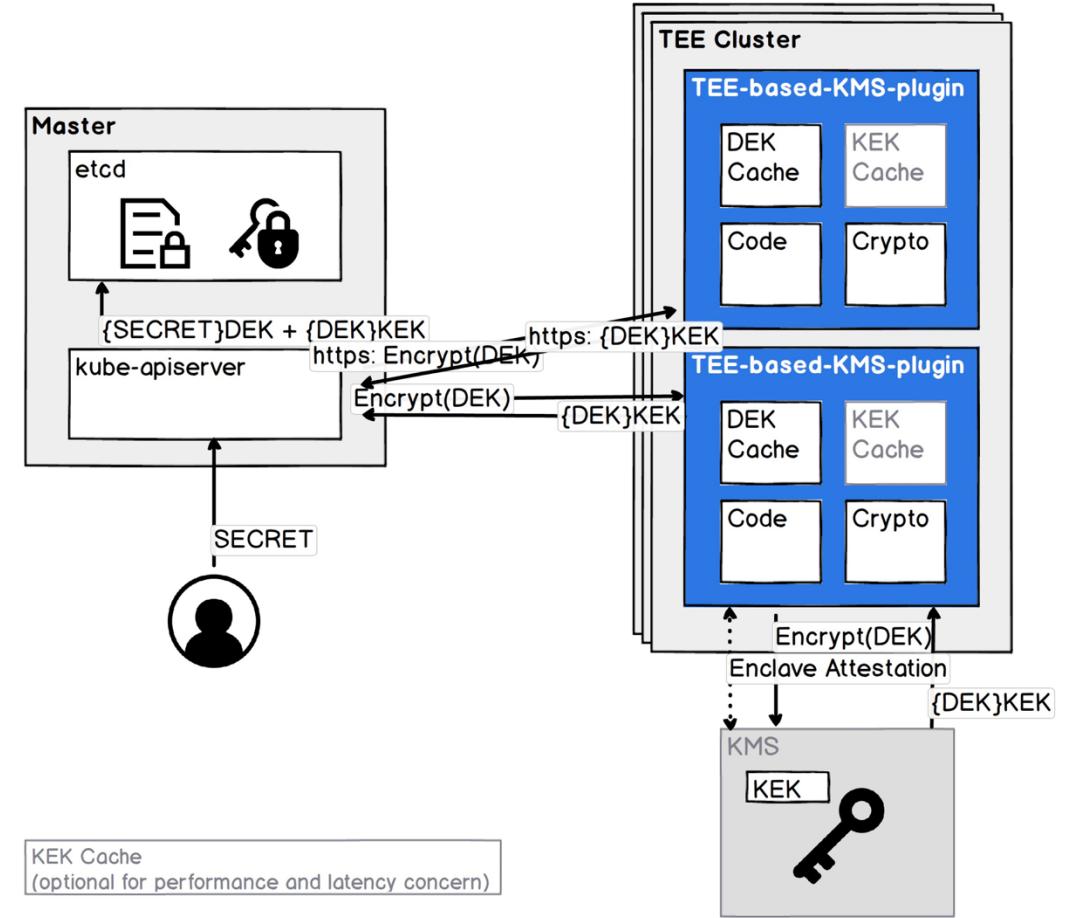
# KMS Plugin as a Service

## Motivation

- SGX physical servers do not meet API servers' performance requirements

## Solution

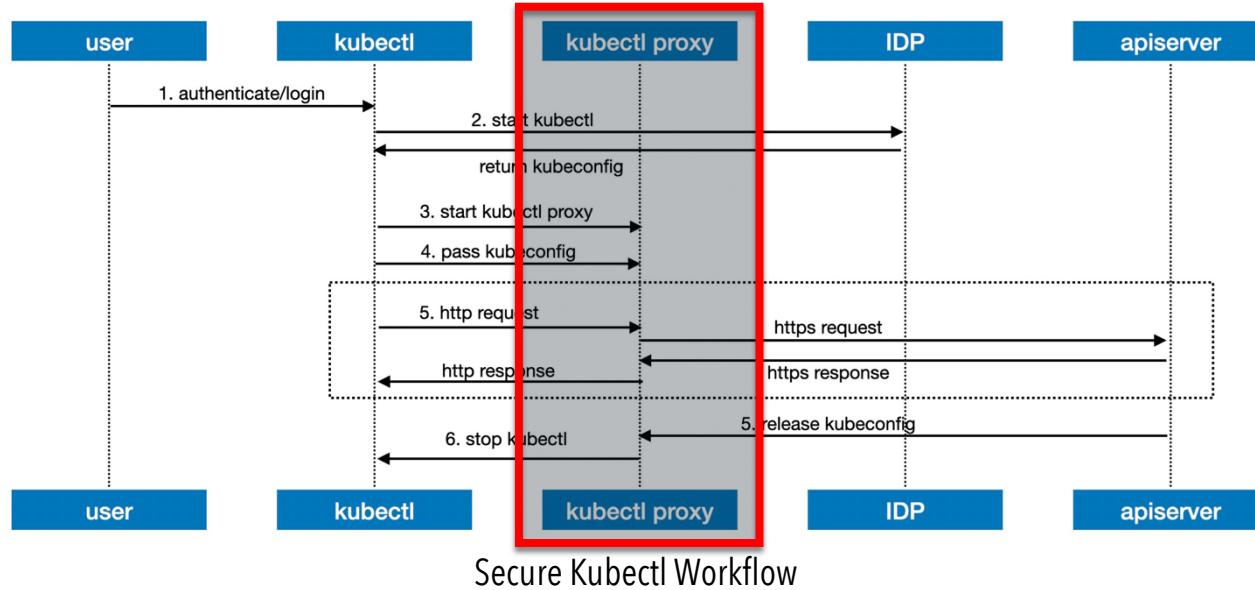
- Same TEE-based KMS-plugin runtime
- Deployment modes
  - N ( $\geq 3$ ) SGX servers deployed w/ sgx-device-plugin daemonset [1]
  - kms-plugins deployed as deployment
- Interfaces
  - https + connection reuse
  - certificate: similar to apiserver  $\Leftrightarrow$  etcd (X.509)
- Version-based key synchronization
- Adaption
  - apiserver KMS provider endpoint to support https endpoint
  - KMS plugin to support https



[1] <https://github.com/AliyunContainerService/sgx-device-plugin>



# Secure Kubectl



- **Design goal**

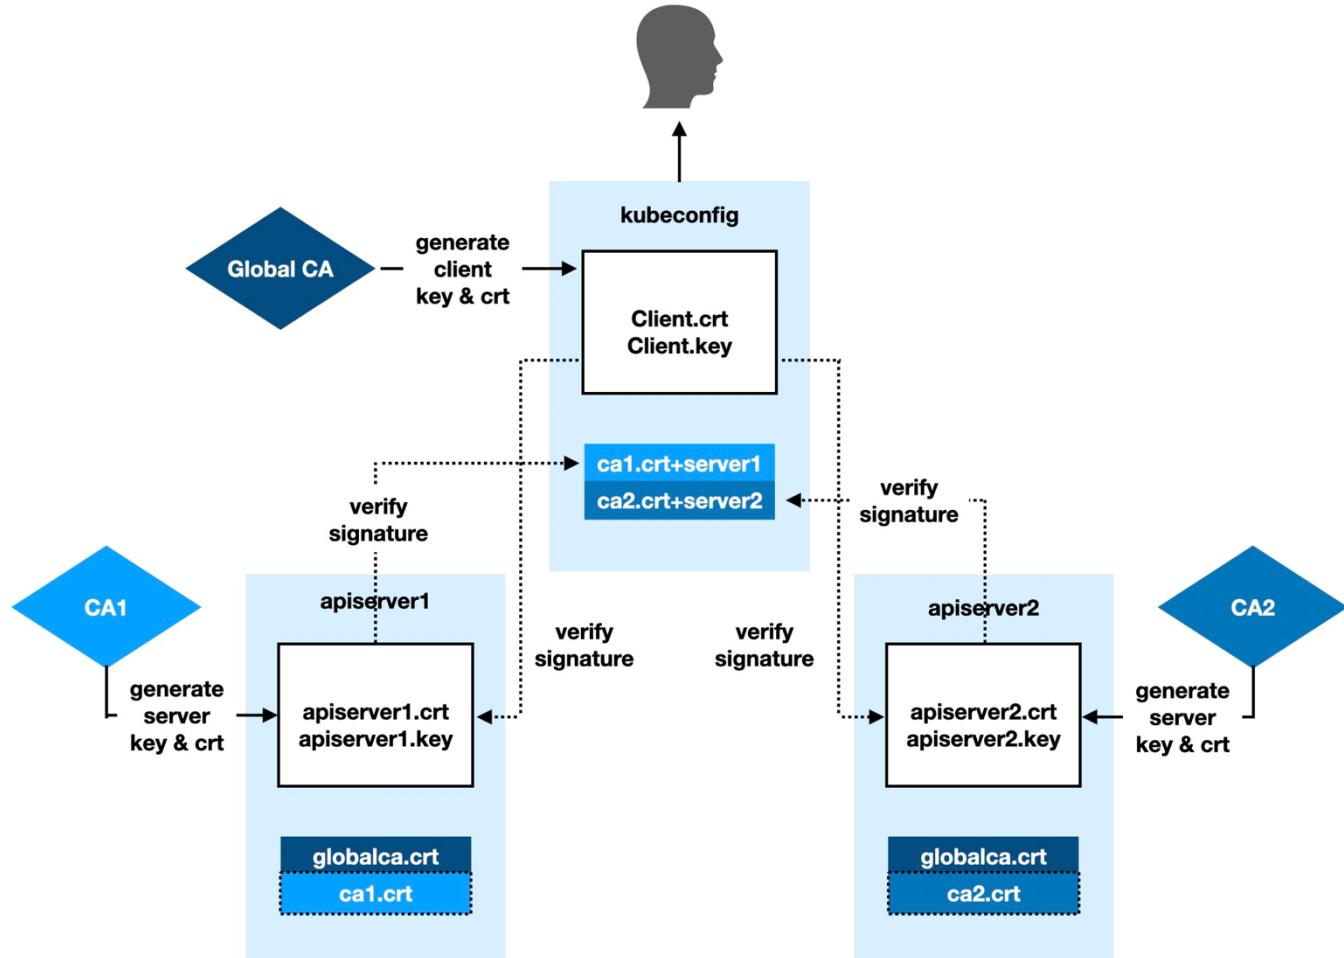
- kubeconfig transparent to kubectl users
- kubeconfig credentials binding w/ identity
- kubeconfig only in memory
- TEE as an option

- **Solution**

- Get kubeconfig
  - Relay server mode
  - Non-relay server mode
- Keep kubeconfig in memory
  - proxy mode
  - kubectl ⇄ http/uds ⇄ proxy ⇄ https ⇄ apiserver
- X.509 or OIDC Token



# Secure Kubectl (cont.)



- Global CA
- One kubeconfig for multiple clusters



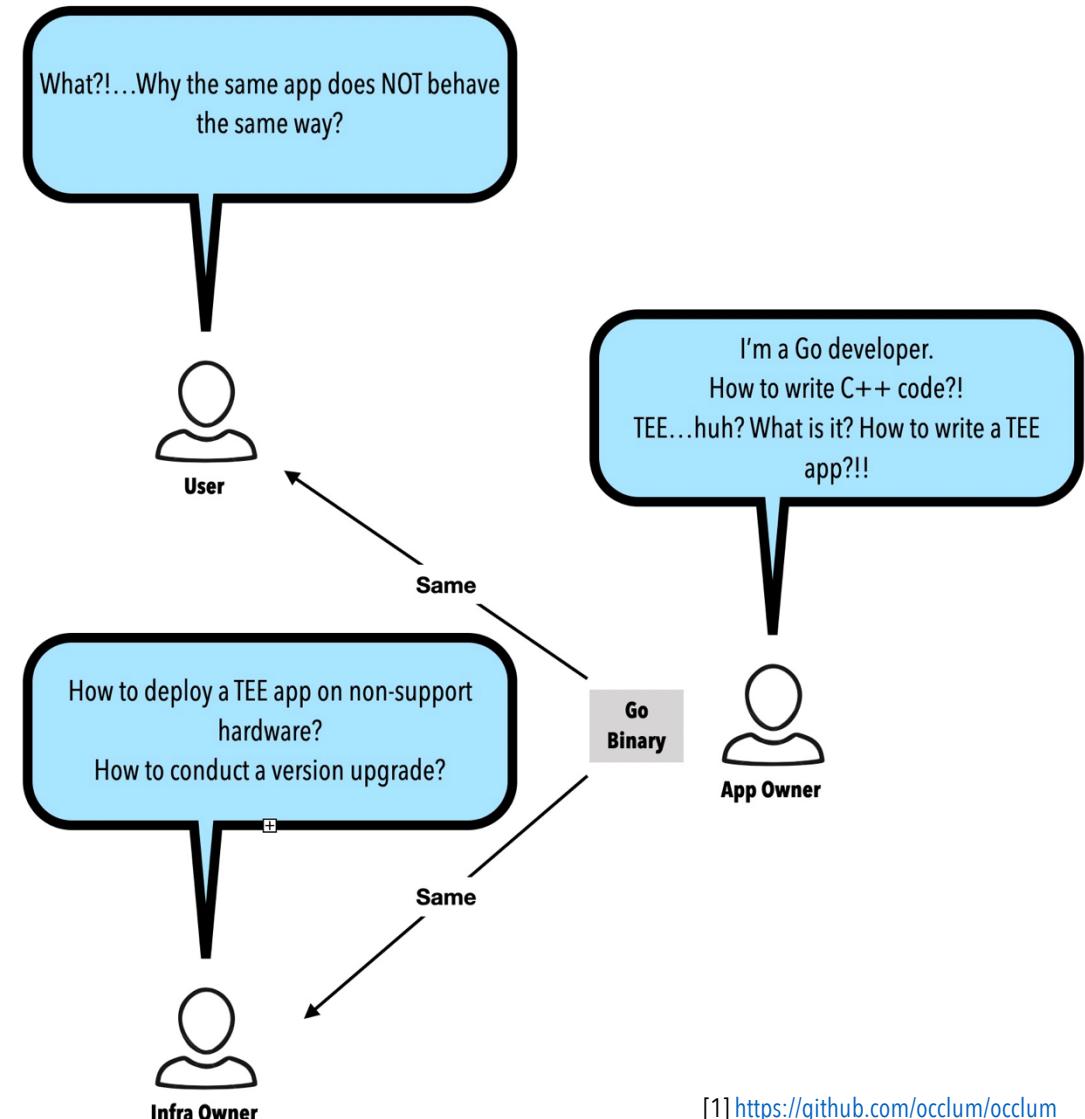
# One binary: TEE Transparency

## • Motivation

- Leverage the same code base, thus the same
  - APIs, logic, iteration plan for developers
  - Experience for users/operators
- TEE as an option, en/disable based on
  - Hardware configurations
  - Biz scenarios

## • Solutions evaluated

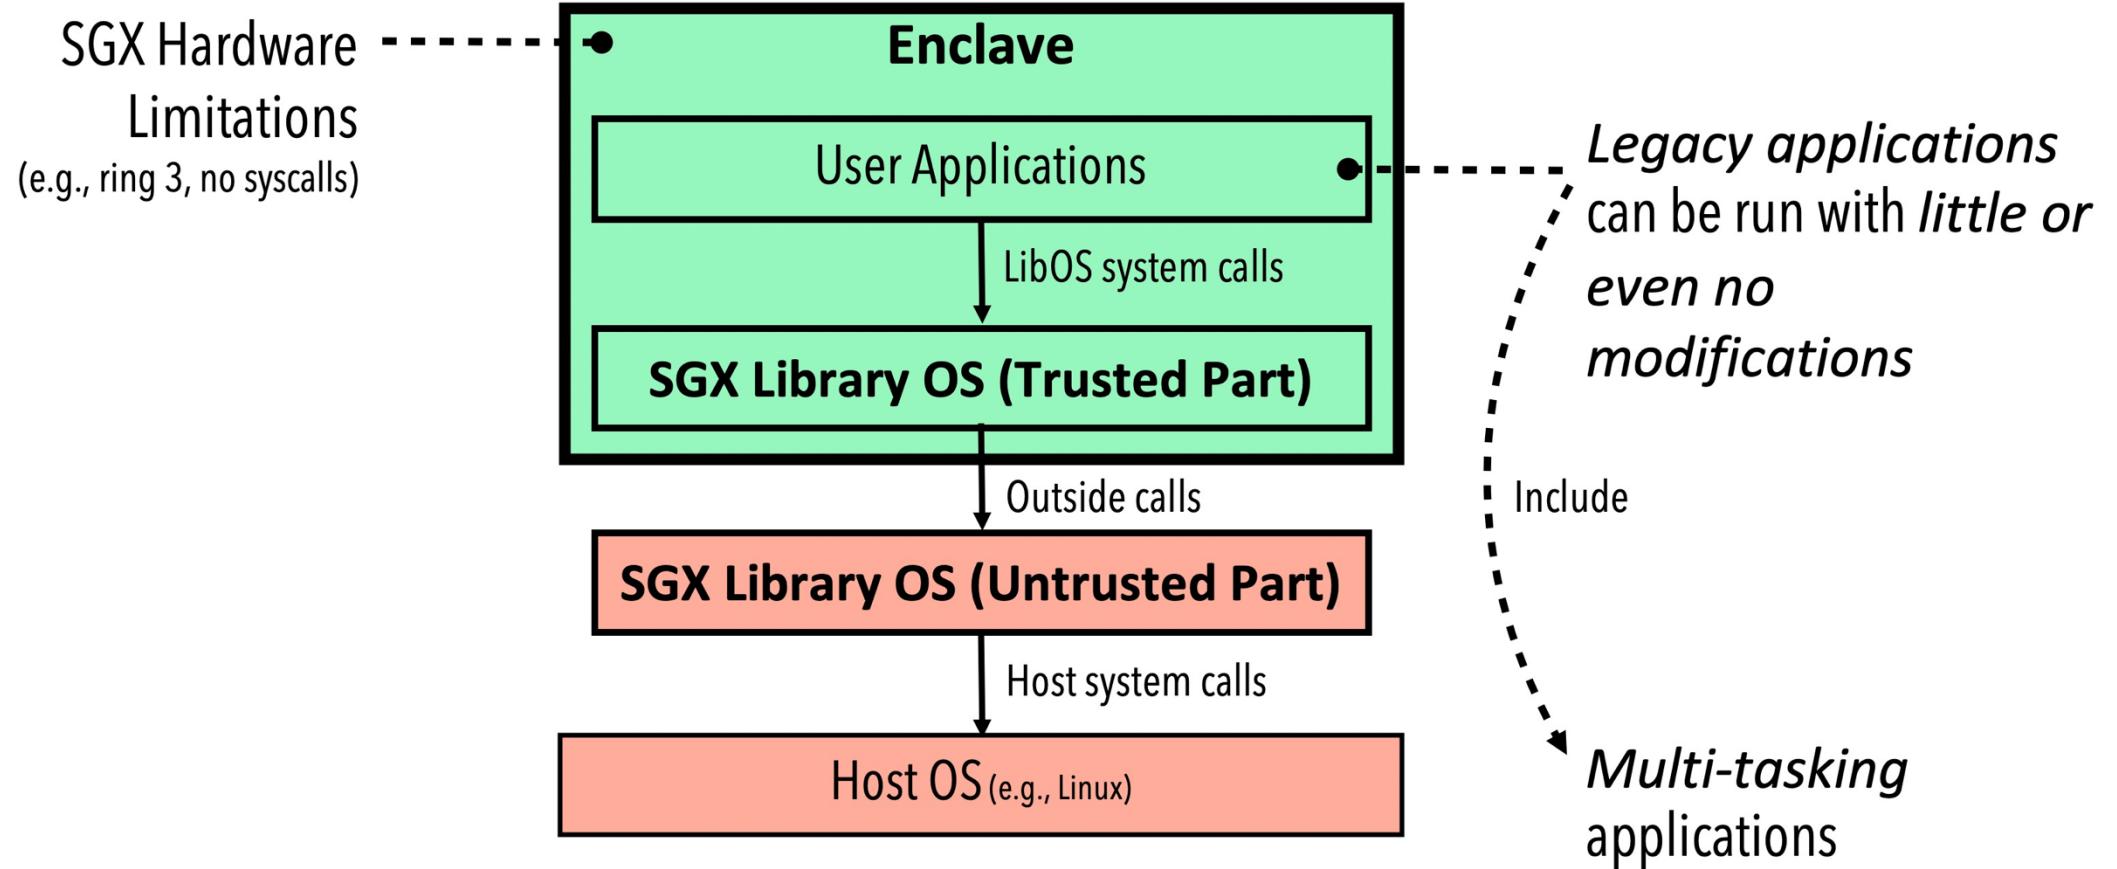
- Go KMS Plugin + C Go SGX Functions
- C++ KMS Plugin
- Rust KMS Plugin
- ✓ Go KMS Plugin + TEE LibOS (Occlum [1])



[1] <https://github.com/occlum/occlum>



# Occlum: SGX Dev Made Easy





# Occlum: Major Features

Usability

**Container-Inspired Interface**

Performance

**Efficient Multitasking**

Functionality

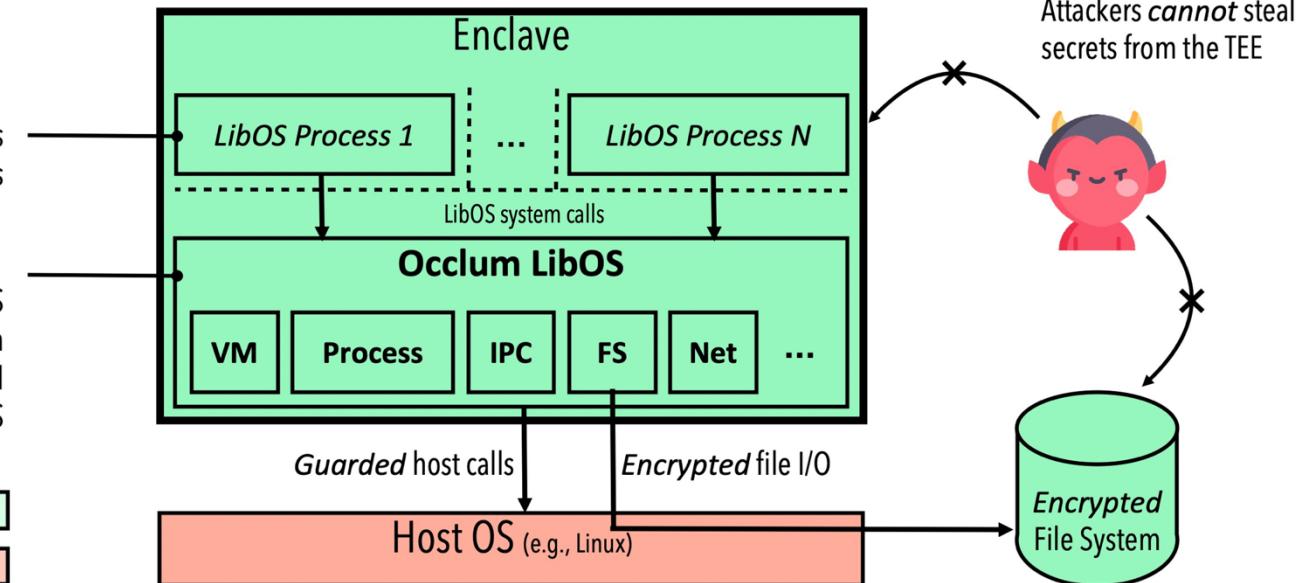
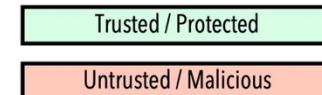
**Multiple File Systems**

Security

**Memory Safety**

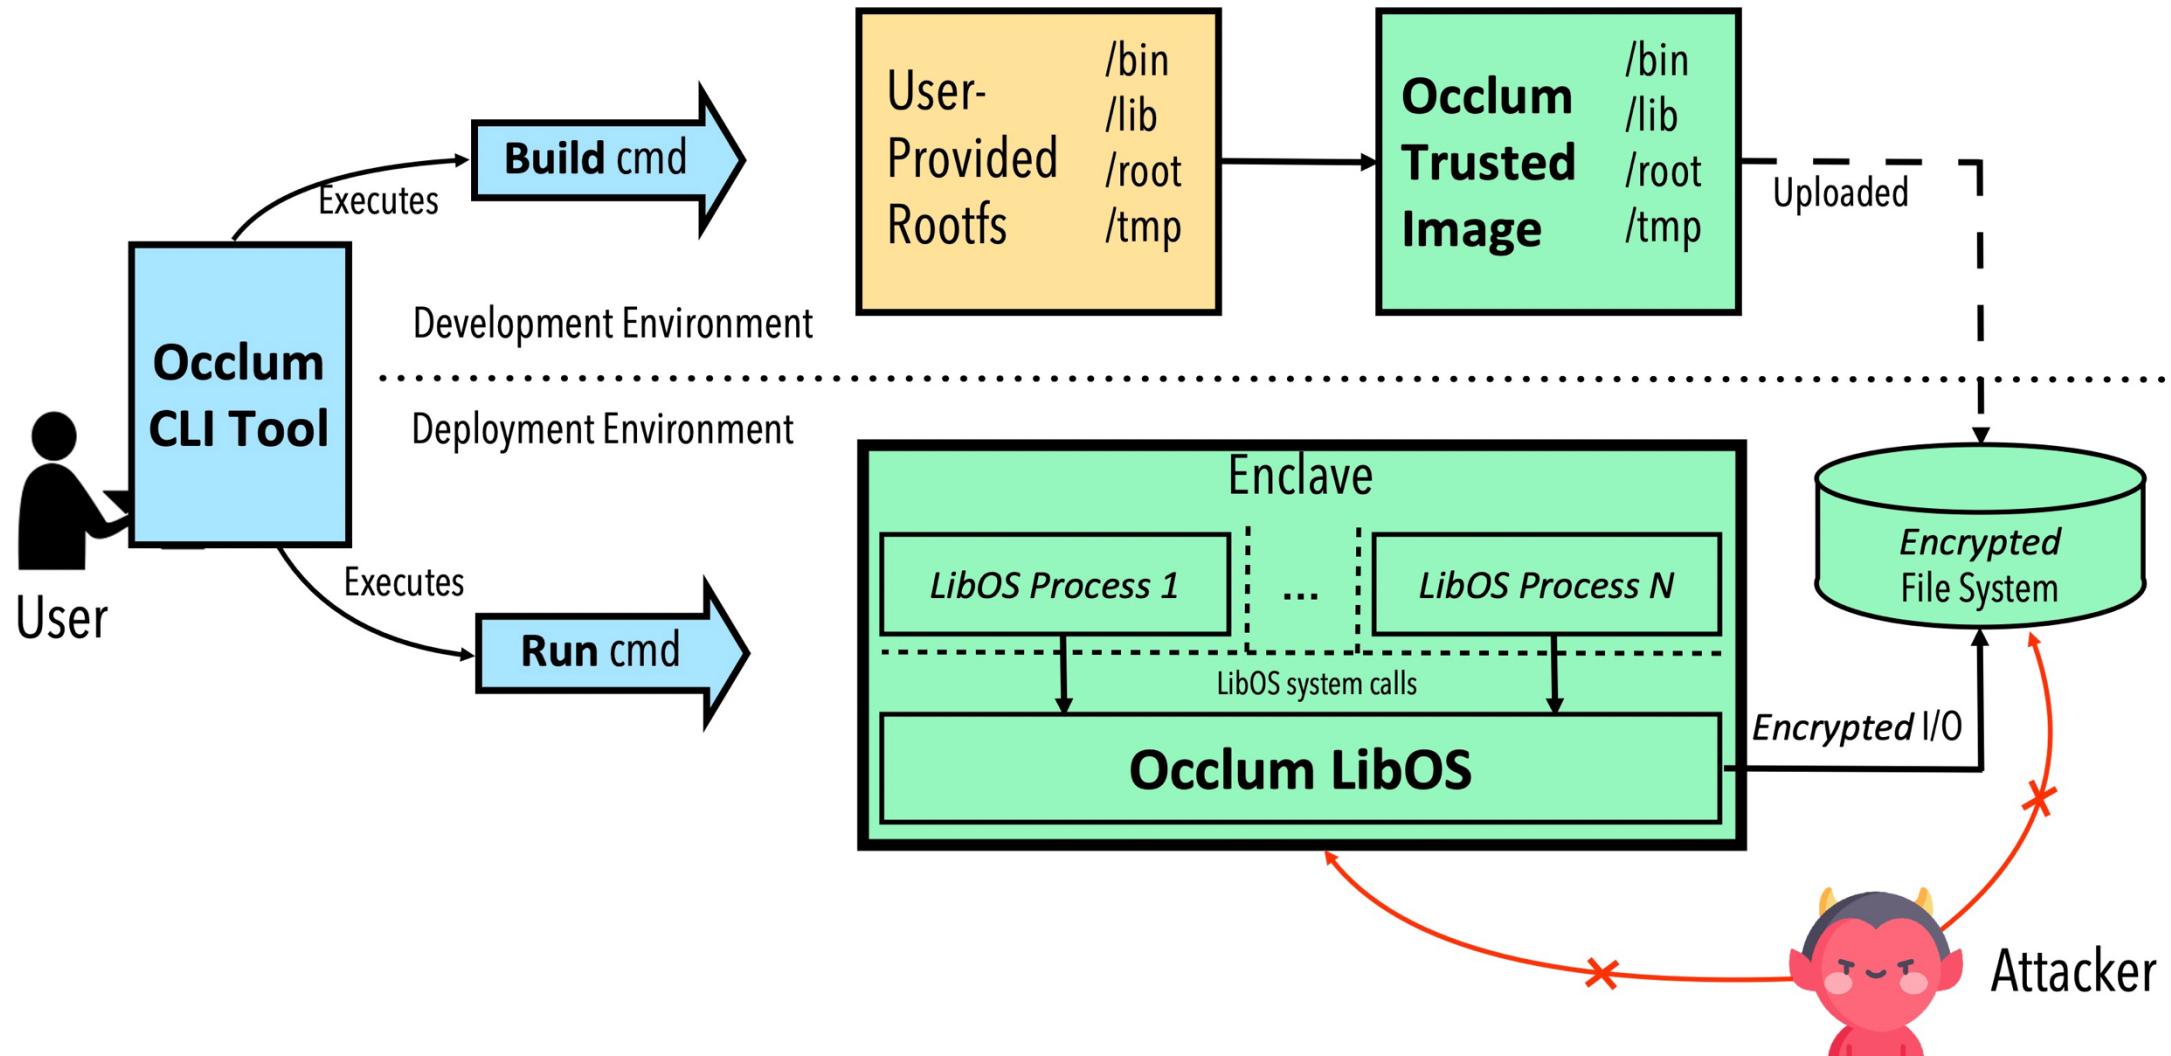
The trusted apps  
requires *minimal* modifications

The LibOS  
is a bridge between  
the trusted apps and  
the untrusted host OS





# Occlum: Container-Inspired Interface





CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*

# Demo



# Demo

- The purpose of this demo is to
  - Demonstrate TEE Transparency w/ Occlum's Golang support
  - Showcase the confidentiality guaranteed by TEE



# Demo

```
iTerm2 Shell Edit View Session Scripts Profiles Toolbelt Window Help
root@a95dbf629742: ~/demos/kms-plugin (ssh)
root@a95dbf629742: ~/demos/kms-plugin (ssh)
kailun.qkl@h07e03158:~$
```

(reverse-i-search)` ':



CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*

# Summary & Plan



CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*

# Summary & Next Steps

- **Summary**
  - A TEE-based E2E solution aiming to guard K8s secrets while in use, at rest, and in transit
  - TEE transparency empowered by LibOS (e.g. Occlum)
- **Where we are?**
  - TEE-based KMS plugin on the way to production
  - TEE-based KMS plugin as a Service - production in plan
  - Secure kubectl already in production; TEE-based secure kubectl completed PoC
  - API server and the rest changes for the TEE-based - PoC stage
- **Where we go?**
  - To keep the production practice
  - To explore more potential intersections between Cloud-Native and Confidential Computing
  - To submit KEPs to upstream for related changes



# Special Thanks

We'd like to thank the contributors to this work:

Yinyong Zheng, Wei Zhang, Yong Zhang, Jun Chen, Yang Yang, Ai Jing, Huanxin Shi, Xiaoyun Mao, Ziteng Xu, Junxian Xiao, Shoumeng Yan and the whole Occlum Team



CLOUD NATIVE + OPEN SOURCE

*Virtual Summit China 2020*